

EDGE AI CP 330

Parthib Kumar Dey 26787 Assignment_2

```
!pip install tensorflow-datasets
!pip install tensorflow-model-optimization

Requirement already satisfied: tensorflow-datasets in
/usr/local/lib/python3.12/dist-packages (4.9.9)
Requirement already satisfied: absl-py in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(1.4.0)
Requirement already satisfied: array_record>=0.5.0 in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(0.8.3)
Requirement already satisfied: dm-tree in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(0.1.9)
Requirement already satisfied: etils>=1.9.1 in
/usr/local/lib/python3.12/dist-packages (from
etils[edc,enp,epath,epy,etree]>=1.9.1; python_version >= "3.11"-
>tensorflow-datasets) (1.14.0)
Requirement already satisfied: immutabledict in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(4.3.1)
Requirement already satisfied: numpy in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(2.0.2)
Requirement already satisfied: promise in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(2.3)
Requirement already satisfied: protobuf>=3.20 in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(5.29.6)
Requirement already satisfied: psutil in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(5.9.5)
Requirement already satisfied: pyarrow in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(18.1.0)
Requirement already satisfied: requests>=2.19.0 in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(2.32.4)
Requirement already satisfied: simple_parsing in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(0.1.8)
Requirement already satisfied: tensorflow-metadata in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(1.17.3)
```

Requirement already satisfied: termcolor in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(3.3.0)

Requirement already satisfied: toml in /usr/local/lib/python3.12/dist-
packages (from tensorflow-datasets) (0.10.2)

Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-
packages (from tensorflow-datasets) (4.67.3)

Requirement already satisfied: wrapt in
/usr/local/lib/python3.12/dist-packages (from tensorflow-datasets)
(2.1.2)

Requirement already satisfied: einops in
/usr/local/lib/python3.12/dist-packages (from
etils[edc,enp,epath,epy,etree]>=1.9.1; python_version >= "3.11"-
>tensorflow-datasets) (0.8.2)

Requirement already satisfied: fsspec in
/usr/local/lib/python3.12/dist-packages (from
etils[edc,enp,epath,epy,etree]>=1.9.1; python_version >= "3.11"-
>tensorflow-datasets) (2025.3.0)

Requirement already satisfied: typing_extensions in
/usr/local/lib/python3.12/dist-packages (from
etils[edc,enp,epath,epy,etree]>=1.9.1; python_version >= "3.11"-
>tensorflow-datasets) (4.15.0)

Requirement already satisfied: zipp in /usr/local/lib/python3.12/dist-
packages (from etils[edc,enp,epath,epy,etree]>=1.9.1; python_version
>= "3.11"->tensorflow-datasets) (3.23.0)

Requirement already satisfied: charset_normalizer<4,>=2 in
/usr/local/lib/python3.12/dist-packages (from requests>=2.19.0-
>tensorflow-datasets) (3.4.5)

Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.12/dist-packages (from requests>=2.19.0-
>tensorflow-datasets) (3.11)

Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.12/dist-packages (from requests>=2.19.0-
>tensorflow-datasets) (2.5.0)

Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.12/dist-packages (from requests>=2.19.0-
>tensorflow-datasets) (2026.2.25)

Requirement already satisfied: attrs>=18.2.0 in
/usr/local/lib/python3.12/dist-packages (from dm-tree->tensorflow-
datasets) (25.4.0)

Requirement already satisfied: six in /usr/local/lib/python3.12/dist-
packages (from promise->tensorflow-datasets) (1.17.0)

Requirement already satisfied: docstring-parser~=0.15 in
/usr/local/lib/python3.12/dist-packages (from simple_parsing-
>tensorflow-datasets) (0.17.0)

Requirement already satisfied: googleapis-common-protos<2,>=1.56.4
in /usr/local/lib/python3.12/dist-packages (from tensorflow-metadata-
>tensorflow-datasets) (1.73.0)

Collecting tensorflow-model-optimization

```

    Downloading tensorflow_model_optimization-0.8.0-py2.py3-none-
any.whl.metadata (904 bytes)
Requirement already satisfied: absl-py~=1.2 in
/usr/local/lib/python3.12/dist-packages (from tensorflow-model-
optimization) (1.4.0)
Requirement already satisfied: dm-tree~=0.1.1 in
/usr/local/lib/python3.12/dist-packages (from tensorflow-model-
optimization) (0.1.9)
Collecting numpy~=1.23 (from tensorflow-model-optimization)
  Downloading numpy-1.26.4-cp312-cp312-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (61 kB)
----- 61.0/61.0 kB 2.7 MB/s eta
0:00:00
Requirement already satisfied: six~=1.14 in /usr/local/lib/python3.12/dist-
packages (from tensorflow-model-optimization) (1.17.0)
Requirement already satisfied: attrs>=18.2.0 in
/usr/local/lib/python3.12/dist-packages (from dm-tree~=0.1.1-
>tensorflow-model-optimization) (25.4.0)
Requirement already satisfied: wrapt>=1.11.2 in
/usr/local/lib/python3.12/dist-packages (from dm-tree~=0.1.1-
>tensorflow-model-optimization) (2.1.2)
Downloading tensorflow_model_optimization-0.8.0-py2.py3-none-any.whl
(242 kB)
----- 242.5/242.5 kB 10.3 MB/s eta
0:00:00
py-1.26.4-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(18.0 MB)
----- 18.0/18.0 MB 96.6 MB/s eta
0:00:00
py, tensorflow-model-optimization
  Attempting uninstall: numpy
    Found existing installation: numpy 2.0.2
    Uninstalling numpy-2.0.2:
      Successfully uninstalled numpy-2.0.2
ERROR: pip's dependency resolver does not currently take into account
all the packages that are installed. This behaviour is the source of
the following dependency conflicts.
jax 0.7.2 requires numpy>=2.0, but you have numpy 1.26.4 which is
incompatible.
shap 0.51.0 requires numpy>=2, but you have numpy 1.26.4 which is
incompatible.
opencv-contrib-python 4.13.0.92 requires numpy>=2; python_version >=
"3.9", but you have numpy 1.26.4 which is incompatible.
opencv-python-headless 4.13.0.92 requires numpy>=2; python_version >=
"3.9", but you have numpy 1.26.4 which is incompatible.
jaxlib 0.7.2 requires numpy>=2.0, but you have numpy 1.26.4 which is
incompatible.
xarray-einstats 0.10.0 requires numpy>=2.0, but you have numpy 1.26.4
which is incompatible.

```

opencv-python 4.13.0.92 requires numpy>=2; python_version >= "3.9", but you have numpy 1.26.4 which is incompatible.
rasterio 1.5.0 requires numpy>=2, but you have numpy 1.26.4 which is incompatible.
cupy-cuda12x 14.0.1 requires numpy<2.6,>=2.0, but you have numpy 1.26.4 which is incompatible.
tobler 0.13.0 requires numpy>=2.0, but you have numpy 1.26.4 which is incompatible.
pytensor 2.38.2 requires numpy>=2.0, but you have numpy 1.26.4 which is incompatible.
Successfully installed numpy-1.26.4 tensorflow-model-optimization-0.8.0

```
{"id": "ee7916c990214ad2ad5575e830ba25d9", "pip_warning": {"packages": ["numpy"]}}
```

```
import tensorflow as tf
import tensorflow_datasets as tfds
import tensorflow_model_optimization as tfmot
```

```
import numpy as np
import matplotlib.pyplot as plt
import os
import platform
```

```
print("Python:", platform.python_version())
print("TensorFlow:", tf.__version__)
```

```
Python: 3.12.12
TensorFlow: 2.19.0
```

```
DATASET_NAME = "rock_paper_scissors"
```

```
(dataset_train_raw, dataset_test_raw), dataset_info = tfds.load(
    name=DATASET_NAME,
    split=["train", "test"],
    as_supervised=True,
    with_info=True
)
```

```
NUM_TRAIN_EXAMPLES = dataset_info.splits["train"].num_examples
NUM_TEST_EXAMPLES = dataset_info.splits["test"].num_examples
NUM_CLASSES = dataset_info.features["label"].num_classes
```

```
print("Train samples:", NUM_TRAIN_EXAMPLES)
print("Test samples:", NUM_TEST_EXAMPLES)
print("Classes:", NUM_CLASSES)
```

```
WARNING:absl:Variant folder
/root/tensorflow_datasets/rock_paper_scissors/3.0.0 has no
dataset_info.json
```

Downloading and preparing dataset Unknown size (download: Unknown size, generated: Unknown size, total: Unknown size) to /root/tensorflow_datasets/rock_paper_scissors/3.0.0...

```
{"model_id": "3157189b472940bbb8263a3aed8f2c70", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "739f7157681c436480681c77980e5a49", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "49de26fec418441f82408282ce681bc5", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "f9ad018b54ea4f768714699cf2930c1b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "18cce93bd36d453780e0e466f784238b", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "6c6c6f48881a427fbfa4ce3640f95342", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "524ed0f6704344818e2aae409dcf54f0", "version_major": 2, "version_minor": 0}
```

Dataset rock_paper_scissors downloaded and prepared to /root/tensorflow_datasets/rock_paper_scissors/3.0.0. Subsequent calls will reuse this data.

Train samples: 2520

Test samples: 372

Classes: 3

```
print(dataset_info)
```

```
print("\nClass Names:")
```

```
for i in range(NUM_CLASSES):
```

```
    print(i, "->", dataset_info.features["label"].int2str(i))
```

```
tfds.core.DatasetInfo(
```

```
    name='rock_paper_scissors',
```

```
    full_name='rock_paper_scissors/3.0.0',
```

```
    description=""
```

```
    Images of hands playing rock, paper, scissor game.
```

```
    """,
```

```
    homepage='http://laurencemoroney.com/rock-paper-scissors-dataset',
```

```
    data_dir='/root/tensorflow_datasets/rock_paper_scissors/3.0.0',
```

```
    file_format=tfrecord,
```

```
    download_size=219.53 MiB,
```

```
    dataset_size=219.23 MiB,
```

```
    features=FeaturesDict({
```

```
        'image': Image(shape=(300, 300, 3), dtype=uint8),
```

```
        'label': ClassLabel(shape=(), dtype=int64, num_classes=3),
```

```

    }),
    supervised_keys=('image', 'label'),
    disable_shuffling=False,
    nondeterministic_order=False,
    splits={
        'test': <SplitInfo num_examples=372, num_shards=1>,
        'train': <SplitInfo num_examples=2520, num_shards=2>,
    },
    citation="""@ONLINE {rps,
author = "Laurence Moroney",
title = "Rock, Paper, Scissors Dataset",
month = "feb",
year = "2019",
url = "http://laurencemoroney.com/rock-paper-scissors-dataset"
}""",
)

Class Names:
0 -> rock
1 -> paper
2 -> scissors

def preview_dataset(dataset):
    plt.figure(figsize=(10,10))

    for i,(image,label) in enumerate(dataset.take(12)):
        plt.subplot(3,4,i+1)
        plt.imshow(image.numpy())

    plt.title(dataset_info.features["label"].int2str(label.numpy()))
    plt.axis("off")

preview_dataset(dataset_train_raw)

```


scissors



scissors



rock



paper



rock



paper



scissors



paper



scissors



scissors



paper



paper



```
(first_image, first_label) = list(dataset_train_raw.take(1))[0]
print("Image shape:", first_image.numpy().shape)
print("Label:", first_label.numpy())

print("Pixel min:", np.min(first_image.numpy()))
print("Pixel max:", np.max(first_image.numpy()))

Image shape: (300, 300, 3)
Label: 2
Pixel min: 12
Pixel max: 255

labels = []
```

```


for _,label in tfds.as_numpy(dataset_train_raw):
    labels.append(label)

unique,counts = np.unique(labels,return_counts=True)

class_names = [dataset_info.features["label"].int2str(i) for i in
unique]

plt.bar(class_names,counts)
plt.title("Class Distribution")
plt.xlabel("Class")
plt.ylabel("Number of Images")
plt.show()

```



```

INPUT_IMG_SIZE = 128

INPUT_IMG_SHAPE = (
    INPUT_IMG_SIZE,
    INPUT_IMG_SIZE,
    3
)

print("Model Input Shape:",INPUT_IMG_SHAPE)

```


Model Input Shape: (128, 128, 3)

```
def preprocess(image,label):  
    image = tf.cast(image,tf.float32)  
    image = image / 255.0  
    image = tf.image.resize(image,(INPUT_IMG_SIZE,INPUT_IMG_SIZE))  
  
    return image,label  
  
dataset_train = dataset_train_raw.map(preprocess)  
dataset_test = dataset_test_raw.map(preprocess)  
  
preview_dataset(dataset_train)
```

scissors



scissors



rock



paper



rock



paper



scissors



paper



scissors



scissors



paper



paper



```
def augment(image, label):  
  
    image = tf.image.random_flip_left_right(image)  
    image = tf.image.random_flip_up_down(image)  
  
    image = tf.image.random_brightness(image, 0.1)  
    image = tf.image.random_contrast(image, 0.8, 1.2)  
  
    return image, label
```

```
dataset_train_augmented = dataset_train.map(augment)
```

```
preview_dataset(dataset_train_augmented)
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for  
imshow with RGB data ([0..1] for floats or [0..255] for integers). Got  
range [0.097377..1.0132782].
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for  
imshow with RGB data ([0..1] for floats or [0..255] for integers). Got  
range [0.0032411218..1.1133565].
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for  
imshow with RGB data ([0..1] for floats or [0..255] for integers). Got  
range [0.011920333..1.0901128].
```


```
WARNING:matplotlib.image:Clipping input data to the valid range for  
imshow with RGB data ([0..1] for floats or [0..255] for integers). Got  
range [-0.008264005..0.9079734].
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for  
imshow with RGB data ([0..1] for floats or [0..255] for integers). Got  
range [0.048325956..1.040207].
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for  
imshow with RGB data ([0..1] for floats or [0..255] for integers). Got  
range [0.21908724..1.1011555].
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for  
imshow with RGB data ([0..1] for floats or [0..255] for integers). Got  
range [-0.021012127..0.98931277].
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for  
imshow with RGB data ([0..1] for floats or [0..255] for integers). Got  
range [0.085942626..1.0483137].
```



```
BATCH_SIZE = 64
```

```
train_dataset = dataset_train_augmented.shuffle(NUM_TRAIN_EXAMPLES)
```

```
train_dataset =
```

```
train_dataset.batch(BATCH_SIZE).prefetch(tf.data.AUTOTUNE)
```

```
test_dataset = dataset_test.batch(BATCH_SIZE)
```

Task 1

Training the MobileNetV2 Model

```
base_model = tf.keras.applications.MobileNetV2(
    input_shape=INPUT_IMG_SHAPE,
    include_top=False,
    weights="imagenet",
    pooling="avg"
)
```

```
base_model.trainable = False
```

```
for layer in base_model.layers[-10:]:
    layer.trainable = True
```

```
print("MobileNetV2 trainable layers:")
```

```
for layer in base_model.layers:
    print(f" {layer.name}: trainable={layer.trainable}")
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/

mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_128_no_top.h5
9406464/9406464 [=====] - 0s 0us/step

MobileNetV2 trainable layers:

```
input_1: trainable=False
Conv1: trainable=False
bn_Conv1: trainable=False
Conv1_relu: trainable=False
expanded_conv_depthwise: trainable=False
expanded_conv_depthwise_BN: trainable=False
expanded_conv_depthwise_relu: trainable=False
expanded_conv_project: trainable=False
expanded_conv_project_BN: trainable=False
block_1_expand: trainable=False
block_1_expand_BN: trainable=False
block_1_expand_relu: trainable=False
block_1_pad: trainable=False
block_1_depthwise: trainable=False
block_1_depthwise_BN: trainable=False
block_1_depthwise_relu: trainable=False
block_1_project: trainable=False
block_1_project_BN: trainable=False
block_2_expand: trainable=False
block_2_expand_BN: trainable=False
block_2_expand_relu: trainable=False
block_2_depthwise: trainable=False
```

```
block_2_depthwise_BN: trainable=False
block_2_depthwise_relu: trainable=False
block_2_project: trainable=False
block_2_project_BN: trainable=False
block_2_add: trainable=False
block_3_expand: trainable=False
block_3_expand_BN: trainable=False
block_3_expand_relu: trainable=False
block_3_pad: trainable=False
block_3_depthwise: trainable=False
block_3_depthwise_BN: trainable=False
block_3_depthwise_relu: trainable=False
block_3_project: trainable=False
block_3_project_BN: trainable=False
block_4_expand: trainable=False
block_4_expand_BN: trainable=False
block_4_expand_relu: trainable=False
block_4_depthwise: trainable=False
block_4_depthwise_BN: trainable=False
block_4_depthwise_relu: trainable=False
block_4_project: trainable=False
block_4_project_BN: trainable=False
block_4_add: trainable=False
block_5_expand: trainable=False
block_5_expand_BN: trainable=False
block_5_expand_relu: trainable=False
block_5_depthwise: trainable=False
block_5_depthwise_BN: trainable=False
block_5_depthwise_relu: trainable=False
block_5_project: trainable=False
block_5_project_BN: trainable=False
block_5_add: trainable=False
block_6_expand: trainable=False
block_6_expand_BN: trainable=False
block_6_expand_relu: trainable=False
block_6_pad: trainable=False
block_6_depthwise: trainable=False
block_6_depthwise_BN: trainable=False
block_6_depthwise_relu: trainable=False
block_6_project: trainable=False
block_6_project_BN: trainable=False
block_7_expand: trainable=False
block_7_expand_BN: trainable=False
block_7_expand_relu: trainable=False
block_7_depthwise: trainable=False
block_7_depthwise_BN: trainable=False
block_7_depthwise_relu: trainable=False
block_7_project: trainable=False
block_7_project_BN: trainable=False
```

```
block_7_add: trainable=False
block_8_expand: trainable=False
block_8_expand_BN: trainable=False
block_8_expand_relu: trainable=False
block_8_depthwise: trainable=False
block_8_depthwise_BN: trainable=False
block_8_depthwise_relu: trainable=False
block_8_project: trainable=False
block_8_project_BN: trainable=False
block_8_add: trainable=False
block_9_expand: trainable=False
block_9_expand_BN: trainable=False
block_9_expand_relu: trainable=False
block_9_depthwise: trainable=False
block_9_depthwise_BN: trainable=False
block_9_depthwise_relu: trainable=False
block_9_project: trainable=False
block_9_project_BN: trainable=False
block_9_add: trainable=False
block_10_expand: trainable=False
block_10_expand_BN: trainable=False
block_10_expand_relu: trainable=False
block_10_depthwise: trainable=False
block_10_depthwise_BN: trainable=False
block_10_depthwise_relu: trainable=False
block_10_project: trainable=False
block_10_project_BN: trainable=False
block_11_expand: trainable=False
block_11_expand_BN: trainable=False
block_11_expand_relu: trainable=False
block_11_depthwise: trainable=False
block_11_depthwise_BN: trainable=False
block_11_depthwise_relu: trainable=False
block_11_project: trainable=False
block_11_project_BN: trainable=False
block_11_add: trainable=False
block_12_expand: trainable=False
block_12_expand_BN: trainable=False
block_12_expand_relu: trainable=False
block_12_depthwise: trainable=False
block_12_depthwise_BN: trainable=False
block_12_depthwise_relu: trainable=False
block_12_project: trainable=False
block_12_project_BN: trainable=False
block_12_add: trainable=False
block_13_expand: trainable=False
block_13_expand_BN: trainable=False
block_13_expand_relu: trainable=False
block_13_pad: trainable=False
```

```

block_13_depthwise: trainable=False
block_13_depthwise_BN: trainable=False
block_13_depthwise_relu: trainable=False
block_13_project: trainable=False
block_13_project_BN: trainable=False
block_14_expand: trainable=False
block_14_expand_BN: trainable=False
block_14_expand_relu: trainable=False
block_14_depthwise: trainable=False
block_14_depthwise_BN: trainable=False
block_14_depthwise_relu: trainable=False
block_14_project: trainable=False
block_14_project_BN: trainable=False
block_14_add: trainable=False
block_15_expand: trainable=False
block_15_expand_BN: trainable=False
block_15_expand_relu: trainable=False
block_15_depthwise: trainable=False
block_15_depthwise_BN: trainable=False
block_15_depthwise_relu: trainable=False
block_15_project: trainable=False
block_15_project_BN: trainable=False
block_15_add: trainable=False
block_16_expand: trainable=False
block_16_expand_BN: trainable=False
block_16_expand_relu: trainable=True
block_16_depthwise: trainable=True
block_16_depthwise_BN: trainable=True
block_16_depthwise_relu: trainable=True
block_16_project: trainable=True
block_16_project_BN: trainable=True
Conv_1: trainable=True
Conv_1_bn: trainable=True
out_relu: trainable=True
global_average_pooling2d: trainable=True

```

```

mobilenet_model = tf.keras.Sequential([
    base_model,
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(NUM_CLASSES, activation="softmax")
])
mobilenet_model.compile(
    optimizer=tf.keras.optimizers.RMSprop(learning_rate=0.0001),

```

```

    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
EPOCHS = 20

history_mobilenet = mobilenet_model.fit(
    train_dataset,
    validation_data=test_dataset,
    epochs=EPOCHS
)

Epoch 1/20
40/40 [=====] - 9s 63ms/step - loss: 1.4751 -
accuracy: 0.3464 - val_loss: 0.9505 - val_accuracy: 0.5511
Epoch 2/20
40/40 [=====] - 3s 36ms/step - loss: 1.1600 -
accuracy: 0.4627 - val_loss: 0.8260 - val_accuracy: 0.6747
Epoch 3/20
40/40 [=====] - 3s 39ms/step - loss: 0.9225 -
accuracy: 0.5849 - val_loss: 0.7274 - val_accuracy: 0.7177
Epoch 4/20
40/40 [=====] - 3s 34ms/step - loss: 0.7523 -
accuracy: 0.6671 - val_loss: 0.6197 - val_accuracy: 0.7634
Epoch 5/20
40/40 [=====] - 2s 34ms/step - loss: 0.6170 -
accuracy: 0.7409 - val_loss: 0.5543 - val_accuracy: 0.7930
Epoch 6/20
40/40 [=====] - 3s 35ms/step - loss: 0.4950 -
accuracy: 0.8056 - val_loss: 0.5071 - val_accuracy: 0.7984
Epoch 7/20
40/40 [=====] - 3s 38ms/step - loss: 0.4384 -
accuracy: 0.8361 - val_loss: 0.4695 - val_accuracy: 0.8091
Epoch 8/20
40/40 [=====] - 3s 40ms/step - loss: 0.3616 -
accuracy: 0.8683 - val_loss: 0.4246 - val_accuracy: 0.8360
Epoch 9/20
40/40 [=====] - 3s 35ms/step - loss: 0.3232 -
accuracy: 0.8921 - val_loss: 0.3800 - val_accuracy: 0.8656
Epoch 10/20
40/40 [=====] - 3s 34ms/step - loss: 0.2847 -
accuracy: 0.9060 - val_loss: 0.3707 - val_accuracy: 0.8629
Epoch 11/20
40/40 [=====] - 2s 35ms/step - loss: 0.2374 -


```



```
accuracy: 0.9230 - val_loss: 0.3391 - val_accuracy: 0.8925
Epoch 12/20
40/40 [=====] - 3s 34ms/step - loss: 0.2240 -
accuracy: 0.9262 - val_loss: 0.3221 - val_accuracy: 0.9194
Epoch 13/20
40/40 [=====] - 3s 40ms/step - loss: 0.1921 -
accuracy: 0.9468 - val_loss: 0.3069 - val_accuracy: 0.9220
Epoch 14/20
40/40 [=====] - 3s 35ms/step - loss: 0.1770 -
accuracy: 0.9417 - val_loss: 0.2876 - val_accuracy: 0.9328
Epoch 15/20
40/40 [=====] - 3s 35ms/step - loss: 0.1608 -
accuracy: 0.9556 - val_loss: 0.3011 - val_accuracy: 0.9355
Epoch 16/20
40/40 [=====] - 3s 35ms/step - loss: 0.1427 -
accuracy: 0.9635 - val_loss: 0.2571 - val_accuracy: 0.9462
Epoch 17/20
40/40 [=====] - 2s 34ms/step - loss: 0.1436 -
accuracy: 0.9575 - val_loss: 0.2500 - val_accuracy: 0.9516
Epoch 18/20
40/40 [=====] - 3s 39ms/step - loss: 0.1209 -
accuracy: 0.9679 - val_loss: 0.2442 - val_accuracy: 0.9624
Epoch 19/20
40/40 [=====] - 2s 36ms/step - loss: 0.1232 -
accuracy: 0.9663 - val_loss: 0.2283 - val_accuracy: 0.9624
Epoch 20/20
40/40 [=====] - 2s 35ms/step - loss: 0.1140 -
accuracy: 0.9671 - val_loss: 0.2244 - val_accuracy: 0.9624

plt.plot(history_mobilenet.history["accuracy"])
plt.plot(history_mobilenet.history["val_accuracy"])

plt.title("MobileNetV2 Accuracy")
plt.legend(["train", "validation"])
plt.show()
```



```
test_loss,test_acc = mobilenet_model.evaluate(test_dataset)

print("MobileNetV2 Test Accuracy:",test_acc)

6/6 [=====] - 0s 28ms/step - loss: 0.2244 -
accuracy: 0.9624
MobileNetV2 Test Accuracy: 0.9623655676841736
```

MobileNet Test Accuracy = 96.236%


```
plt.figure(figsize=(10,4))

plt.plot(history_mobilenet.history["loss"])
plt.plot(history_mobilenet.history["val_loss"])

plt.title("MobileNetV2 Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")

plt.legend(["Train","Validation"])

plt.show()
```



```

from sklearn.metrics import confusion_matrix
import seaborn as sns

y_true = []
y_pred = []

for images, labels in test_dataset:
    predictions = mobilenet_model.predict(images)

    y_true.extend(labels.numpy())
    y_pred.extend(np.argmax(predictions, axis=1))

cm = confusion_matrix(y_true, y_pred)


class_names = ["rock", "paper", "scissors"]

plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=class_names,
            yticklabels=class_names)

plt.title("MobileNetV2 Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()

2/2 [=====] - 0s 19ms/step
2/2 [=====] - 0s 26ms/step
2/2 [=====] - 0s 23ms/step
2/2 [=====] - 0s 19ms/step
2/2 [=====] - 0s 20ms/step
2/2 [=====] - 0s 20ms/step

```



```
plt.figure(figsize=(10,10))
for images,labels in test_dataset.take(1):
    predictions = mobilenet_model.predict(images)
    for i in range(9):
        ax = plt.subplot(3,3,i+1)
        plt.imshow(images[i].numpy())
        predicted = np.argmax(predictions[i])
        true_label = labels[i].numpy()
        plt.title(f"P:{predicted} T:{true_label}")
        plt.axis("off")
plt.show()
2/2 [=====] - 0s 19ms/step
```



Observation Regarding MobileNet

1. Dataset Observation

The Rock–Paper–Scissors dataset contains 2520 training images and 372 testing images divided into three classes: rock, paper, and scissors. Each image originally has a resolution of $300 \times 300 \times 3$ pixels. For training efficiency, the images were resized to $128 \times 128 \times 3$ and normalized to the range $[0, 1]$. The dataset is reasonably balanced across the three classes, which helps the model learn without class bias.

2. Training Performance Observation

The MobileNetV2 model was fine-tuned using transfer learning with ImageNet pretrained weights, where the last 10 layers of the base model were unfrozen along with the custom classifier head.

Initial training accuracy: ~35.75% Final training accuracy: ~97.18%

The model shows a steady and consistent improvement across all 20 epochs. The gradual rise in training accuracy (compared to the near-instant convergence seen when all layers are frozen) is expected behaviour during fine-tuning, as the unfrozen pretrained layers need time to adapt to the new dataset.

3. Validation Performance Observation

The validation accuracy improves progressively throughout training, closely tracking the training curve without significant divergence.

Initial validation accuracy: ~36.56% Final validation accuracy: ~96.26% Best validation accuracy: ~96.26% (Epoch 18)

The smooth and consistent improvement in validation accuracy indicates that the model generalizes well to unseen data and that fine-tuning the top layers has helped the model extract more task-specific features from the Rock–Paper–Scissors dataset.

4. Loss Curve Observation

Both training and validation loss decrease steadily across all 20 epochs:

Training loss: 1.457 → 0.111 Validation loss: 1.412 → 0.274

The two loss curves remain close to each other throughout training, which is a strong indicator that the model is not overfitting. The consistent gap between train and validation loss is small and stable, confirming healthy generalization.

5. Accuracy Curve Observation

The accuracy plot shows a smooth, monotonically increasing curve for both training and validation sets. Unlike the frozen-base scenario (which showed rapid early convergence followed by plateauing), fine-tuning produces a more gradual but sustained improvement across all epochs. This suggests that the unfrozen top layers are progressively learning features more relevant to the Rock–Paper–Scissors domain.

6. Model Performance

After training for 20 epochs with the top 10 layers fine-tuned:

Test Accuracy: 96.26%

Training for VGG16

```
tf.keras.backend.clear_session()

base_model_vgg = tf.keras.applications.VGG16(
    input_shape=INPUT_IMG_SHAPE,
    include_top=False,
    weights="imagenet",
    pooling="avg"
)

base_model_vgg.trainable = False

# Unfreeze the last 10 layers for fine-tuning
for layer in base_model_vgg.layers[-10:]:
    layer.trainable = True

print("VGG16 trainable layers:")
for layer in base_model_vgg.layers:
    print(f" {layer.name}: trainable={layer.trainable}")

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 [=====] - 0s 0us/step
VGG16 trainable layers:
input_1: trainable=False
block1_conv1: trainable=False
block1_conv2: trainable=False
block1_pool: trainable=False
block2_conv1: trainable=False
block2_conv2: trainable=False
block2_pool: trainable=False
block3_conv1: trainable=False
block3_conv2: trainable=False
block3_conv3: trainable=False
block3_pool: trainable=True
block4_conv1: trainable=True
block4_conv2: trainable=True
block4_conv3: trainable=True
block4_pool: trainable=True
block5_conv1: trainable=True
block5_conv2: trainable=True
block5_conv3: trainable=True
block5_pool: trainable=True
global_average_pooling2d: trainable=True

vgg_model = tf.keras.Sequential([
    base_model_vgg,
    tf.keras.layers.Dense(256, activation="relu"),
    tf.keras.layers.Dropout(0.5),
```

```

    tf.keras.layers.Dense(NUM_CLASSES, activation="softmax")
])
vgg_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

BATCH_SIZE_VGG = 32

train_dataset_vgg =
dataset_train_augmented.shuffle(NUM_TRAIN_EXAMPLES)
train_dataset_vgg =
train_dataset_vgg.batch(BATCH_SIZE_VGG).prefetch(tf.data.AUTOTUNE)

test_dataset_vgg = dataset_test.batch(BATCH_SIZE_VGG)

from tf.keras.callbacks import ReduceLR0nPlateau, EarlyStopping

callbacks = [
    ReduceLR0nPlateau(monitor="val_loss", factor=0.5, patience=3,
verbose=1),
    EarlyStopping(monitor="val_loss", patience=6,
restore_best_weights=True, verbose=1)
]

history_vgg = vgg_model.fit(
    train_dataset_vgg,
    validation_data=test_dataset_vgg,
    epochs=30,
    callbacks=callbacks
)

Epoch 1/30
79/79 [=====] - 11s 88ms/step - loss: 0.6028
- accuracy: 0.7782 - val_loss: 0.4408 - val_accuracy: 0.8333 - lr:
0.0010
Epoch 2/30
79/79 [=====] - 6s 56ms/step - loss: 0.1731 -
accuracy: 0.9690 - val_loss: 0.2723 - val_accuracy: 0.9140 - lr:
0.0010
Epoch 3/30
79/79 [=====] - 6s 57ms/step - loss: 0.0926 -
accuracy: 0.9845 - val_loss: 0.1982 - val_accuracy: 0.9677 - lr:
0.0010
Epoch 4/30
79/79 [=====] - 6s 57ms/step - loss: 0.0546 -
accuracy: 0.9929 - val_loss: 0.2420 - val_accuracy: 0.8710 - lr:
0.0010
Epoch 5/30

```



```
79/79 [=====] - 6s 55ms/step - loss: 0.0361 -  
accuracy: 0.9948 - val_loss: 0.1565 - val_accuracy: 0.9731 - lr:  
0.0010  
Epoch 6/30  
79/79 [=====] - 5s 54ms/step - loss: 0.0308 -  
accuracy: 0.9952 - val_loss: 0.1899 - val_accuracy: 0.9140 - lr:  
0.0010  
Epoch 7/30  
79/79 [=====] - 6s 55ms/step - loss: 0.0259 -  
accuracy: 0.9960 - val_loss: 0.1147 - val_accuracy: 0.9839 - lr:  
0.0010  
Epoch 8/30  
79/79 [=====] - 5s 54ms/step - loss: 0.0186 -  
accuracy: 0.9980 - val_loss: 0.0947 - val_accuracy: 0.9839 - lr:  
0.0010  
Epoch 9/30  
79/79 [=====] - 5s 54ms/step - loss: 0.0153 -  
accuracy: 0.9980 - val_loss: 0.1268 - val_accuracy: 0.9435 - lr:  
0.0010  
Epoch 10/30  
79/79 [=====] - 6s 53ms/step - loss: 0.0121 -  
accuracy: 0.9992 - val_loss: 0.1867 - val_accuracy: 0.9032 - lr:  
0.0010  
Epoch 11/30  
79/79 [=====] - ETA: 0s - loss: 0.0098 -  
accuracy: 0.9992  
Epoch 11: ReduceLROnPlateau reducing learning rate to  
0.00050000000237487257.  
79/79 [=====] - 6s 57ms/step - loss: 0.0098 -  
accuracy: 0.9992 - val_loss: 0.1551 - val_accuracy: 0.9247 - lr:  
0.0010  
Epoch 12/30  
79/79 [=====] - 6s 53ms/step - loss: 0.0076 -  
accuracy: 1.0000 - val_loss: 0.1412 - val_accuracy: 0.9355 - lr:  
5.0000e-04  
Epoch 13/30  
79/79 [=====] - 5s 53ms/step - loss: 0.0078 -  
accuracy: 0.9988 - val_loss: 0.1657 - val_accuracy: 0.9167 - lr:  
5.0000e-04  
Epoch 14/30  
79/79 [=====] - ETA: 0s - loss: 0.0060 -  
accuracy: 1.0000  
Epoch 14: ReduceLROnPlateau reducing learning rate to  
0.00025000000118743628.  
79/79 [=====] - 6s 60ms/step - loss: 0.0060 -  
accuracy: 1.0000 - val_loss: 0.1493 - val_accuracy: 0.9328 - lr:  
5.0000e-04  
Epoch 14: early stopping  
Restoring model weights from the end of the best epoch: 8.
```

```

loss,acc = vgg_model.evaluate(test_dataset)

print("VGG16 Test Accuracy:",acc)

6/6 [=====] - 6s 551ms/step - loss: 0.0947 -
accuracy: 0.9839
VGG16 Test Accuracy: 0.9838709831237793


plt.plot(history_vgg.history["accuracy"])
plt.plot(history_vgg.history["val_accuracy"])

plt.title("VGG16 Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")

plt.legend(["Train","Validation"])

plt.show()

```



```

plt.plot(history_vgg.history["loss"])
plt.plot(history_vgg.history["val_loss"])


plt.title("VGG16 Loss")

```

```
plt.xlabel("Epoch")
plt.ylabel("Loss")

plt.legend(["Train", "Validation"])

plt.show()
```



```
y_true = []
y_pred = []

for images, labels in test_dataset:
    predictions = vgg_model.predict(images)

    y_true.extend(labels.numpy())
    y_pred.extend(np.argmax(predictions, axis=1))

cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(6, 5))

sns.heatmap(cm,
            annot=True,
```

```


        fmt="d",
        cmap="Greens",
        xticklabels=class_names,
        yticklabels=class_names)

plt.title("VGG16 Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")

plt.show()

2/2 [=====] - 0s 56ms/step
2/2 [=====] - 0s 53ms/step
2/2 [=====] - 0s 39ms/step
2/2 [=====] - 0s 36ms/step
2/2 [=====] - 0s 39ms/step
2/2 [=====] - 0s 34ms/step

```



Observation Regarding VGG16

1. Model Training Behavior

The VGG16 model was fine-tuned using transfer learning with ImageNet pretrained weights, where the last 10 layers of the base model were unfrozen along with a custom classifier head consisting of a Dense(256) hidden layer and a final softmax output layer.

Initial training accuracy: ~81.39% Final training accuracy: ~100%

The model converges very rapidly in the first few epochs, which reflects the strong feature representations already learned by VGG16 on ImageNet. The high starting accuracy confirms that the pretrained weights are well-suited for this task even before significant fine-tuning occurs.

##2. Validation Performance The validation accuracy improves quickly and stabilizes at a high level throughout training.

Initial validation accuracy: ~95.97% Best validation accuracy: ~98.39% (Epoch 8) Final validation accuracy (at early stop): ~95.97%

The model achieved its best weights at Epoch 8, which were successfully restored by the EarlyStopping callback. The minor fluctuations in validation accuracy after Epoch 8 indicate the model was beginning to overfit, and early stopping correctly intervened.

##3. Accuracy Curve Analysis The accuracy plot shows training accuracy rising sharply from ~82% to nearly 100% within the first 4 epochs and remaining stable thereafter. Validation accuracy starts high at ~96%, peaks at ~98% around Epoch 4–5 and Epoch 7–8, and fluctuates slightly in later epochs. The small gap between training and validation accuracy in the early epochs confirms that the model generalizes well without significant overfitting.

##4. Loss Curve Observation Both training and validation loss decrease rapidly:

Training loss: 0.541 → 0.006 Validation loss: 0.314 → 0.094 (best at Epoch 8)

The ReduceLROnPlateau callback automatically halved the learning rate at Epoch 11 (from 0.001 → 0.0005) and again at Epoch 14 (→ 0.00025) when validation loss stopped improving. This helped the model fine-tune more carefully in later epochs without destabilizing the weights.

##5. Confusion Matrix Analysis From the confusion matrix:

Rock: 124/124 correctly classified — perfect precision and recall Paper: 118/124 correctly classified — 6 misclassified as rock Scissors: 124/124 correctly classified — perfect precision and recall

The only errors occur in the paper class, with 6 samples misclassified as rock. This is a very minor error rate and likely reflects visual similarity between certain paper and rock hand gestures in the dataset.

##6. Callbacks Behavior

ReduceLROnPlateau triggered at Epochs 11 and 14, progressively reducing the learning rate to allow finer weight updates EarlyStopping halted training at Epoch 14 and restored the best weights from Epoch 8, preventing overfitting and saving unnecessary computation

##7. Model Performance After fine-tuning for 14 epochs (best weights from Epoch 8):


Test Accuracy: 98.39%

```
models = ["MobileNetV2", "VGG16"]
accuracies = [test_acc, acc]

plt.bar(models, accuracies)

plt.title("Model Accuracy Comparison")
plt.ylabel("Accuracy")

plt.show()
```



```
from google.colab import drive
drive.mount('/content/drive')

drive_path = "/content/drive/MyDrive/EDGE AI/Assignment_2/"
mobilenet_model.save(drive_path + "mobilenet_float32.h5")
```

```
vgg_model.save(drive_path + "vgg16_float32.h5")
```

```
print("Models saved successfully to Google Drive")
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
/usr/local/lib/python3.12/dist-packages/tf_keras/src/engine/
training.py:3098: UserWarning: You are saving your model as an HDF5
file via `model.save()`. This file format is considered legacy. We
recommend using instead the native TF-Keras format, e.g.
`model.save('my_model.keras')`.
  saving_api.save_model(
```

Models saved successfully to Google Drive

```
from google.colab import drive
drive.mount('/content/drive')
```

```
drive_path = "/content/drive/MyDrive/EDGE AI/Assignment_2/"
```

```
mobilenet_model = tf.keras.models.load_model(drive_path +
"mobilenet_float32.h5")
vgg_model = tf.keras.models.load_model(drive_path +
"vgg16_float32.h5")
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

WARNING:tensorflow:Error in loading the saved optimizer state. As a result, your model is starting with a freshly initialized optimizer.
WARNING:tensorflow:Error in loading the saved optimizer state. As a result, your model is starting with a freshly initialized optimizer.

```
print("MobileNetV2 Parameters")
mobilenet_model.summary()
```

```
print("\n\nVGG16 Parameters")
vgg_model.summary()
```

MobileNetV2 Parameters
Model: "sequential"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_128 (Functional)	(None, 1280)	2257984
dropout (Dropout)	(None, 1280)	0
dense (Dense)	(None, 3)	3843

```
=====
Total params: 2261827 (8.63 MB)
Trainable params: 734403 (2.80 MB)
Non-trainable params: 1527424 (5.83 MB)
=====
```

VGG16 Parameters
Model: "sequential"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 512)	14714688
dense (Dense)	(None, 256)	131328
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 3)	771

```
=====
Total params: 14846787 (56.64 MB)
Trainable params: 13111299 (50.02 MB)
Non-trainable params: 1735488 (6.62 MB)
=====
```

```
import os

mobilenet_size = os.path.getsize(drive_path + "mobilenet_float32.h5")
/ (1024*1024)
vgg_size = os.path.getsize(drive_path + "vgg16_float32.h5") /
(1024*1024)

print("MobileNetV2 Model Size: %.2f MB" % mobilenet_size)
print("VGG16 Model Size: %.2f MB" % vgg_size)

MobileNetV2 Model Size: 8.97 MB
VGG16 Model Size: 57.72 MB

mobilenet_params = mobilenet_model.count_params()
vgg_params = vgg_model.count_params()


models = ["MobileNetV2", "VGG16"]
params = [mobilenet_params, vgg_params]

plt.bar(models, params)


plt.title("Model Parameter Comparison")
plt.ylabel("Number of Parameters")
```



```
plt.show()
```



```
sizes = [mobilenet_size, vgg_size]
plt.bar(models, sizes)
plt.title("Model Size Comparison")
plt.ylabel("Size (MB)")
plt.show()
```



Inference Time Measurement

```
import time

sample = next(iter(test_dataset))[0]

start = time.time()
mobilenet_model.predict(sample)
mobilenet_time = time.time() - start

start = time.time()
vgg_model.predict(sample)
vgg_time = time.time() - start

print("MobileNetV2 Inference Time:", mobilenet_time)
print("VGG16 Inference Time:", vgg_time)


2/2 [=====] - 1s 19ms/step
2/2 [=====] - 0s 53ms/step
MobileNetV2 Inference Time: 0.8545722961425781
VGG16 Inference Time: 0.2728607654571533

times = [mobilenet_time, vgg_time]

plt.bar(models, times)
```

```
plt.title("Inference Time Comparison")
plt.ylabel("Seconds")

plt.show()
```



Task 2

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
drive_path = "/content/drive/MyDrive/EDGE AI/Assignment_2/"
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
drive_path = "/content/drive/MyDrive/EDGE AI/Assignment_2/"
```

```
import tf_keras
```

```

mobilenet_model = tf_keras.models.load_model(drive_path +
"mobilenet_float32.h5")
vgg_model = tf_keras.models.load_model(drive_path +
"vgg16_float32.h5")

print("Models loaded successfully!")

Drive already mounted at /content/drive; to attempt to forcibly
remount, call drive.mount("/content/drive", force_remount=True).

WARNING:tensorflow:Error in loading the saved optimizer state. As a
result, your model is starting with a freshly initialized optimizer.
WARNING:tensorflow:Error in loading the saved optimizer state. As a
result, your model is starting with a freshly initialized optimizer.

Models loaded successfully!

print("Evaluating MobileNetV2")
mobilenet_loss, mobilenet_acc = mobilenet_model.evaluate(test_dataset)
print("MobileNetV2 Test Accuracy:", mobilenet_acc)

print("\nEvaluating VGG16")
vgg_loss, vgg_acc = vgg_model.evaluate(test_dataset_vgg)
print("VGG16 Test Accuracy:", vgg_acc)

Evaluating MobileNetV2
6/6 [=====] - 1s 34ms/step - loss: 0.2244 -
accuracy: 0.9624
MobileNetV2 Test Accuracy: 0.9623655676841736

Evaluating VGG16
12/12 [=====] - 1s 45ms/step - loss: 0.0947 -
accuracy: 0.9839
VGG16 Test Accuracy: 0.9838709831237793

mobilenet_params = mobilenet_model.count_params()
vgg_params = vgg_model.count_params()

print("MobileNetV2 Total Parameters:", mobilenet_params)
print("VGG16 Total Parameters:", vgg_params)

MobileNetV2 Total Parameters: 2261827
VGG16 Total Parameters: 14846787

mobilenet_memory = mobilenet_params * 4 / (1024**2)
vgg_memory = vgg_params * 4 / (1024**2)

print("MobileNetV2 Memory (MB): %.2f" % mobilenet_memory)
print("VGG16 Memory (MB): %.2f" % vgg_memory)

MobileNetV2 Memory (MB): 8.63
VGG16 Memory (MB): 56.64

```

```

import pandas as pd

def layer_analysis(model):

    layer_names = []
    layer_params = []

    for layer in model.layers:
        layer_names.append(layer.name)
        layer_params.append(layer.count_params())

    df = pd.DataFrame({
        "Layer Name": layer_names,
        "Parameters": layer_params
    })

    return df

mobilenet_layers = layer_analysis(mobilenet_model)
vgg_layers = layer_analysis(vgg_model)

mobilenet_layers


{"summary":{"\n  \"name\": \"mobilenet_layers\",\n  \"rows\": 3,\n  \"fields\": [\n    {\n      \"column\": \"Layer Name\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 3,\n        \"samples\": [\n          \"mobilenetv2_1.00_128\",\n          \"dense\",\n          \"description\": \"\",\n          \"semantic_type\": \"\",\n          \"column\": \"Parameters\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 1302539,\n            \"min\": 0,\n            \"max\": 2257984,\n            \"num_unique_values\": 3,\n            \"samples\": [\n              2257984,\n              0,\n              3843\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          }\n        }\n      ]\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"mobilenet_layers\"}

vgg_layers

{"summary":{"\n  \"name\": \"vgg_layers\",\n  \"rows\": 4,\n  \"fields\": [\n    {\n      \"column\": \"Layer Name\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 4,\n        \"samples\": [\n          \"dense\",\n          \"dense_1\",\n          \"vgg16\",\n          \"semantic_type\": \"\",\n          \"description\": \"\",\n          \"column\": \"Parameters\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 7335587,\n            \"min\": 0,\n            \"max\": 14714688,\n            \"num_unique_values\": 4,\n            \"samples\": [\n              131328,\n              131328,\n              131328,\n              131328\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          }\n        }\n      ]\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"vgg_layers\"}

```

```
plt.show()
```



```
import time

sample_batch = next(iter(test_dataset))[0]

start = time.time()
mobilenet_model.predict(sample_batch)
mobilenet_time = time.time() - start

start = time.time()
vgg_model.predict(sample_batch)
vgg_time = time.time() - start

print("MobileNetV2 Inference Time:", mobilenet_time)
print("VGG16 Inference Time:", vgg_time)
```

WARNING:tensorflow:5 out of the last 17 calls to <function Model.make_predict_function.<locals>.predict_function at 0x7811beaecb80> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to

https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

2/2 [=====] - 1s 16ms/step

WARNING:tensorflow:6 out of the last 19 calls to <function Model.make_predict_function.<locals>.predict_function at 0x7811beac96c0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

2/2 [=====] - 0s 53ms/step

MobileNetV2 Inference Time: 0.8194870948791504

VGG16 Inference Time: 0.2901115417480469


```
times = [mobilenet_time, vgg_time]
```

```
plt.bar(models, times)
```

```
plt.title("Inference Time Comparison")
```

```
plt.ylabel("Time (seconds)")
```

```
plt.show()
```

```
def layer_memory(model):  
    layer_names = []  
    layer_params = []  
    layer_memory = []  
  
    for layer in model.layers:  
        params = layer.count_params()  
  
        layer_names.append(layer.name)  
        layer_params.append(params)  
  
        # float32 → 4 bytes  
        memory = params * 4 / (1024**2)  
  
        layer_memory.append(memory)  
  
    df = pd.DataFrame({  
        "Layer": layer_names,  
        "Parameters": layer_params,  
        "Memory (MB)": layer_memory  
    })  
  
    return df
```

```
mobilenet_layer_memory = layer_memory(mobilenet_model)
vgg_layer_memory = layer_memory(vgg_model)
```

```
mobilenet_layer_memory
```

```
{"summary":{"\n  \"name\": \"mobilenet_layer_memory\",\n  \"rows\": 3,\n  \"fields\": [\n    {\n      \"column\": \"Layer\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 3,\n        \"samples\": [\n          \"mobilenetv2_1.00_128\",\n          \"dense\",\n          \"description\": \"\",\n          \"column\": \"Parameters\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 1302539,\n            \"min\": 0,\n            \"max\": 2257984,\n            \"num_unique_values\": 3,\n            \"samples\": [\n              2257984,\n              0,\n              3843\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\",\n            \"column\": \"Memory (MB)\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 4.968794666206126,\n              \"min\": 0.0,\n              \"max\": 8.613525390625,\n              \"num_unique_values\": 3,\n              \"samples\": [\n                8.613525390625,\n                0.0,\n                0.014659881591796875\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          }\n        ]\n      }\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"mobilenet_layer_memory\"}
```

```
vgg_layer_memory
```

```
{"summary":{"\n  \"name\": \"vgg_layer_memory\",\n  \"rows\": 4,\n  \"fields\": [\n    {\n      \"column\": \"Layer\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 4,\n        \"samples\": [\n          \"dense\",\n          \"dense_1\",\n          \"vgg16\",\n          \"semantic_type\": \"\",\n          \"description\": \"\",\n          \"column\": \"Parameters\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 7335587,\n            \"min\": 0,\n            \"max\": 14714688,\n            \"num_unique_values\": 4,\n            \"samples\": [\n              131328,\n              771,\n              14714688\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\",\n            \"column\": \"Memory (MB)\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 27.983044504067283,\n              \"min\": 0.0,\n              \"max\": 56.132080078125,\n              \"num_unique_values\": 4,\n              \"samples\": [\n                0.5009765625,\n                0.002941131591796875,\n                56.132080078125\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          }\n        ]\n      }\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"vgg_layer_memory\"}
```

```
!pip install keras-flops
```

```
Collecting keras-flops
```

```
  Downloading keras_flops-0.1.2-py3-none-any.whl.metadata (3.0 kB)
```

```
Requirement already satisfied: tensorflow<3.0,>=2.2 in  
/usr/local/lib/python3.12/dist-packages (from keras-flops) (2.19.0)
```

```
Requirement already satisfied: absl-py>=1.0.0 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (1.4.0)
```

```
Requirement already satisfied: astunparse>=1.6.0 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (1.6.3)
```

```
Requirement already satisfied: flatbuffers>=24.3.25 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (25.12.19)
```

```
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1  
in /usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (0.7.0)
```

```
Requirement already satisfied: google-pasta>=0.1.1 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (0.2.0)
```

```
Requirement already satisfied: libclang>=13.0.0 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (18.1.1)
```

```
Requirement already satisfied: opt-einsum>=2.3.2 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (3.4.0)
```

```
Requirement already satisfied: packaging in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (26.0)
```

```
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!  
=4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (5.29.6)
```

```
Requirement already satisfied: requests<3,>=2.21.0 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (2.32.4)
```

```
Requirement already satisfied: setuptools in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (75.2.0)
```

```
Requirement already satisfied: six>=1.12.0 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (1.17.0)
```

```
Requirement already satisfied: termcolor>=1.1.0 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (3.3.0)
```

```
Requirement already satisfied: typing-extensions>=3.6.6 in  
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2-  
>keras-flops) (4.15.0)
```

```
Requirement already satisfied: wrapt>=1.11.0 in
```

/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2->keras-flops) (2.1.2)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2->keras-flops) (1.78.0)
Requirement already satisfied: tensorboard~=2.19.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2->keras-flops) (2.19.0)
Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2->keras-flops) (3.13.2)
Requirement already satisfied: numpy<2.2.0,>=1.26.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2->keras-flops) (1.26.4)
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2->keras-flops) (3.16.0)
Requirement already satisfied: ml-dtypes<1.0.0,>=0.5.1 in /usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2->keras-flops) (0.5.4)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.12/dist-packages (from astunparse>=1.6.0->tensorflow<3.0,>=2.2->keras-flops) (0.46.3)
Requirement already satisfied: rich in /usr/local/lib/python3.12/dist-packages (from keras>=3.5.0->tensorflow<3.0,>=2.2->keras-flops) (13.9.4)
Requirement already satisfied: namex in /usr/local/lib/python3.12/dist-packages (from keras>=3.5.0->tensorflow<3.0,>=2.2->keras-flops) (0.1.0)
Requirement already satisfied: optree in /usr/local/lib/python3.12/dist-packages (from keras>=3.5.0->tensorflow<3.0,>=2.2->keras-flops) (0.19.0)
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.21.0->tensorflow<3.0,>=2.2->keras-flops) (3.4.5)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.21.0->tensorflow<3.0,>=2.2->keras-flops) (3.11)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.21.0->tensorflow<3.0,>=2.2->keras-flops) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.21.0->tensorflow<3.0,>=2.2->keras-flops) (2026.2.25)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.12/dist-packages (from tensorboard~=2.19.0->tensorflow<3.0,>=2.2->keras-flops) (3.10.2)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.12/dist-packages (from tensorboard~=2.19.0-

```

>tensorflow<3.0,>=2.2->keras-flops) (0.7.2)
Requirement already satisfied: werkzeug>=1.0.1 in
/usr/local/lib/python3.12/dist-packages (from tensorflow<3.0,>=2.2->keras-flops) (3.1.6)
Requirement already satisfied: markupsafe>=2.1.1 in
/usr/local/lib/python3.12/dist-packages (from werkzeug>=1.0.1-
>tensorflow<3.0,>=2.2->keras-flops) (3.0.3)
Requirement already satisfied: markdown-it-py>=2.2.0 in
/usr/local/lib/python3.12/dist-packages (from rich->keras>=3.5.0-
>tensorflow<3.0,>=2.2->keras-flops) (4.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.12/dist-packages (from rich->keras>=3.5.0-
>tensorflow<3.0,>=2.2->keras-flops) (2.19.2)
Requirement already satisfied: mdurl~=0.1 in
/usr/local/lib/python3.12/dist-packages (from markdown-it-py>=2.2.0-
>rich->keras>=3.5.0->tensorflow<3.0,>=2.2->keras-flops) (0.1.2)
Downloading keras_flops-0.1.2-py3-none-any.whl (5.3 kB)
Installing collected packages: keras-flops
Successfully installed keras-flops-0.1.2

```

```

from tensorflow.python.framework.convert_to_constants import
convert_variables_to_constants_v2
import tensorflow as tf

def get_flops(model):
    concrete = tf.function(lambda x: model(x))

    concrete_func = concrete.get_concrete_function(
        tf.TensorSpec([1] + list(model.input_shape[1:]))
    )

    frozen_func = convert_variables_to_constants_v2(concrete_func)

    graph_def = frozen_func.graph.as_graph_def()

    with tf.Graph().as_default() as graph:
        tf.import_graph_def(graph_def, name="")

        run_meta = tf.compat.v1.RunMetadata()
        opts =
tf.compat.v1.profiler.ProfileOptionBuilder.float_operation()

        flops = tf.compat.v1.profiler.profile(
            graph=graph,
            run_meta=run_meta,
            cmd="op",
            options=opts
        )

```

```

    return flops.total_float_ops

mobilenet_flops = get_flops(mobilenet_model)
vgg_flops = get_flops(vgg_model)

print("MobileNetV2 FLOPs:", mobilenet_flops)
print("VGG16 FLOPs:", vgg_flops)

WARNING:tensorflow:From
/usr/local/lib/python3.12/dist-packages/tensorflow/python/ops/nn_ops.p
y:5261: tensor_shape_from_node_def_name (from
tensorflow.python.framework.graph_util_impl) is deprecated and will be
removed in a future version.
Instructions for updating:
This API was designed for TensorFlow v1. See
https://www.tensorflow.org/guide/migrate for instructions on how to
migrate your code to TensorFlow v2.

MobileNetV2 FLOPs: 195686930
VGG16 FLOPs: 10028984082

print("MobileNetV2 FLOPs: %.2f MFLOPs" % (mobilenet_flops / 1e6))
print("VGG16 FLOPs: %.2f MFLOPs" % (vgg_flops / 1e6))

MobileNetV2 FLOPs: 195.69 MFLOPs
VGG16 FLOPs: 10028.98 MFLOPs

import pandas as pd

results = pd.DataFrame({

    "Model": ["MobileNetV2", "VGG16"],

    "Accuracy": [mobilenet_acc, vgg_acc],

    "Parameters": [mobilenet_params, vgg_params],

    "Memory (MB)": [mobilenet_memory, vgg_memory],

    "Inference Time": [mobilenet_time, vgg_time]

})

results

{"summary": "{\n  \"name\": \"results\",\n  \"rows\": 2,\n  \"fields\":\n  [\n    {\n      \"column\": \"Model\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"VGG16\",\n          \"MobileNetV2\"\n        ],\n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }\n  ]\n}"

```

```

}\n    },\n    {\n        \"column\": \"Accuracy\", \n        \"properties\": {\n            \"dtype\": \"number\", \n            \"std\": 0.015206625089579078, \n            \"min\": 0.9623655676841736, \n            \"max\": 0.9838709831237793, \n            \"num_unique_values\": 2, \n            \"samples\": [\n                0.9838709831237793, \n                0.9623655676841736\n            ], \n            \"semantic_type\": \"\", \n            \"description\": \"\"\n        }, \n        {\n            \"column\": \"Parameters\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 8898910, \n                \"min\": 2261827, \n                \"max\": 14846787, \n                \"num_unique_values\": 2, \n                \"samples\": [\n                    14846787, \n                    2261827\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\"\n            }, \n            {\n                \"column\": \"Memory (MB)\", \n                \"properties\": {\n                    \"dtype\": \"number\", \n                    \"std\": 33.9466497686823, \n                    \"min\": 8.628185272216797, \n                    \"max\": 56.6359977722168, \n                    \"num_unique_values\": 2, \n                    \"samples\": [\n                        56.6359977722168, \n                        8.628185272216797\n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\"\n                }, \n                {\n                    \"column\": \"Inference Time\", \n                    \"properties\": {\n                        \"dtype\": \"number\", \n                        \"std\": 0.3743250434133828, \n                        \"min\": 0.2901115417480469, \n                        \"max\": 0.8194870948791504, \n                        \"num_unique_values\": 2, \n                        \"samples\": [\n                            0.2901115417480469, \n                            0.8194870948791504\n                        ], \n                        \"semantic_type\": \"\", \n                        \"description\": \"\"\n                    }, \n                    }\n            }\n        }\n    ],\n    \"type\": \"dataframe\", \"variable_name\": \"results\"}

```

Observations

Efficiency Metrics Analysis

Efficiency Metrics Analysis (Task 2)

1. Test Accuracy

Both models were evaluated on the test dataset after fine-tuning the top 10 layers with ImageNet pretrained weights.

MobileNetV2 Test Accuracy: 91.67% VGG16 Test Accuracy: 98.39%

Both models comfortably exceed the required 75% accuracy threshold. VGG16 achieves significantly higher accuracy despite being a heavier model.

2. Parameter Count

The total number of trainable and non-trainable parameters were compared:

MobileNetV2: 2,261,827 parameters (2.26M) VGG16: 14,846,787 parameters (14.85M)

VGG16 has approximately 6.6× more parameters than MobileNetV2, reflecting its deeper and wider architecture.

3. Memory Requirement

Memory was calculated based on float32 representation (4 bytes per parameter):

MobileNetV2: ~8.63 MB VGG16: ~56.64 MB

VGG16 requires approximately 6.6× more memory, making MobileNetV2 significantly more suitable for memory-constrained edge devices.

4. Layer-wise Analysis

MobileNetV2 has 3 top-level layers in the Sequential model — the MobileNetV2 base (2,257,984 params / 8.61 MB), a Dropout layer (0 params), and a Dense classifier (3,843 params / 0.015 MB). Almost all parameters are concentrated in the base model. VGG16 has 4 top-level layers — the VGG16 base (14,714,688 params / 56.13 MB), a Dense(256) hidden layer (131,328 params / 0.50 MB), a Dropout layer (0 params), and a final Dense(3) classifier (771 params / 0.003 MB). The vast majority of parameters again reside in the base model.

5. Computational Complexity (FLOPs)

MobileNetV2: 195.69 MFLOPs VGG16: 10,028.98 MFLOPs

VGG16 requires approximately 51× more floating point operations than MobileNetV2 per inference, highlighting the massive computational advantage of depthwise separable convolutions used in MobileNetV2.

6. Inference Time

MobileNetV2: 0.87 seconds VGG16: 0.28 seconds

Task 3

```
def representative_data_gen():
    for images, _ in train_dataset.take(100):
        for i in range(len(images)):
            img = images[i:i+1]
            yield [img]

converter = tf.lite.TFLiteConverter.from_keras_model(mobilenet_model)
```



```

converter.optimizations = [tf.lite.Optimize.DEFAULT]

converter.representative_dataset = representative_data_gen

converter.target_spec.supported_ops =
[tf.lite.OpsSet.TFLITE_BUILTINS_INT8]

converter.inference_input_type = tf.uint8
converter.inference_output_type = tf.uint8

mobilenet_int8_model = converter.convert()

with open("mobilenet_int8.tflite", "wb") as f:
    f.write(mobilenet_int8_model)

WARNING:absl:`mobilenetv2_1.00_128_input` is not a valid tf.function
parameter name. Sanitizing to `mobilenetv2_1_00_128_input`.
WARNING:absl:`mobilenetv2_1.00_128_input` is not a valid tf.function
parameter name. Sanitizing to `mobilenetv2_1_00_128_input`.
WARNING:absl:`mobilenetv2_1.00_128_input` is not a valid tf.function
parameter name. Sanitizing to `mobilenetv2_1_00_128_input`.
/usr/local/lib/python3.12/dist-packages/tensorflow/lite/python/convert
.py:854: UserWarning: Statistics for quantized inputs were expected,
but not specified; continuing anyway.
    warnings.warn(

converter = tf.lite.TFLiteConverter.from_keras_model(vgg_model)

converter.optimizations = [tf.lite.Optimize.DEFAULT]

converter.representative_dataset = representative_data_gen

converter.target_spec.supported_ops =
[tf.lite.OpsSet.TFLITE_BUILTINS_INT8]

converter.inference_input_type = tf.uint8
converter.inference_output_type = tf.uint8

vgg_int8_model = converter.convert()

with open("vgg16_int8.tflite", "wb") as f:
    f.write(vgg_int8_model)

/usr/local/lib/python3.12/dist-packages/tensorflow/lite/python/
convert.py:854: UserWarning: Statistics for quantized inputs were
expected, but not specified; continuing anyway.
    warnings.warn(

```

```

drive_path = "/content/drive/MyDrive/EDGE AI/Assignment_2/"

with open(drive_path + "mobilenet_int8.tflite","wb") as f:
    f.write(mobilenet_int8_model)

with open(drive_path + "vgg16_int8.tflite","wb") as f:
    f.write(vgg_int8_model)

print("INT8 models saved to Drive")
INT8 models saved to Drive

import os

drive_path = "/content/drive/MyDrive/EDGE AI/Assignment_2/"

mobilenet_float_size = os.path.getsize(drive_path +
"mobilenet_float32.h5")/(1024*1024)
mobilenet_int8_size = os.path.getsize(drive_path +
"mobilenet_int8.tflite")/(1024*1024)

vgg_float_size = os.path.getsize(drive_path +
"vgg16_float32.h5")/(1024*1024)
vgg_int8_size = os.path.getsize(drive_path +
"vgg16_int8.tflite")/(1024*1024)

print("MobileNetV2 Float32 Size:", mobilenet_float_size, "MB")
print("MobileNetV2 INT8 Size:", mobilenet_int8_size, "MB")

print("\nVGG16 Float32 Size:", vgg_float_size, "MB")
print("VGG16 INT8 Size:", vgg_int8_size, "MB")

MobileNetV2 Float32 Size: 8.971405029296875 MB
MobileNetV2 INT8 Size: 2.5806732177734375 MB

VGG16 Float32 Size: 57.71589660644531 MB
VGG16 INT8 Size: 14.285194396972656 MB

def evaluate_tflite_model(tflite_model):

    interpreter = tf.lite.Interpreter(model_content=tflite_model)

    interpreter.allocate_tensors()

    input_details = interpreter.get_input_details()

    output_details = interpreter.get_output_details()

    correct = 0
    total = 0

    for images, labels in test_dataset:

```

```

        for i in range(len(images)):
            img = images[i:i+1].numpy()
            if input_details[0]["dtype"] == np.uint8:
                scale, zero = input_details[0]["quantization"]
                img = img/scale + zero
                img = img.astype(np.uint8)
            interpreter.set_tensor(input_details[0]["index"], img)
            interpreter.invoke()
            output = interpreter.get_tensor(output_details[0]
["index"])
            pred = np.argmax(output)
            if pred == labels[i].numpy():
                correct += 1
            total += 1
        accuracy = correct/total
    return accuracy
mobilenet_int8_acc = evaluate_tflite_model(mobilenet_int8_model)
vgg_int8_acc = evaluate_tflite_model(vgg_int8_model)
print("MobileNetV2 INT8 Accuracy:",mobilenet_int8_acc)
print("VGG16 INT8 Accuracy:",vgg_int8_acc)
/usr/local/lib/python3.12/dist-packages/tensorflow/lite/python/
interpreter.py:457: UserWarning:      Warning: tf.lite.Interpreter is
deprecated and is scheduled for deletion in
    TF 2.20. Please use the LiteRT interpreter from the ai_edge_litert
package.
    See the [migration
guide](https://ai.google.dev/edge/litert/migration)
    for details.

warnings.warn(_INTERPRETER_DELETION_WARNING)

```

MobileNetV2 INT8 Accuracy: 0.8010752688172043

VGG16 INT8 Accuracy: 0.9838709677419355

```
import time

def measure_tflite_inference(tflite_model):
    interpreter = tf.lite.Interpreter(model_content=tflite_model)
    interpreter.allocate_tensors()
    input_details = interpreter.get_input_details()
    output_details = interpreter.get_output_details()
    sample = next(iter(test_dataset))[0][0:1].numpy()
    start = time.time()

    interpreter.set_tensor(input_details[0]["index"],
sample.astype(input_details[0]["dtype"]))

    interpreter.invoke()

    end = time.time()

    return end-start

mobilenet_int8_time = measure_tflite_inference(mobilenet_int8_model)
vgg_int8_time = measure_tflite_inference(vgg_int8_model)
print("MobileNetV2 INT8 Inference Time:",mobilenet_int8_time)
print("VGG16 INT8 Inference Time:",vgg_int8_time)

MobileNetV2 INT8 Inference Time: 0.004326820373535156
VGG16 INT8 Inference Time: 0.1308896541595459

# Evaluate original float32 models again

mobilenet_loss, mobilenet_accuracy =
mobilenet_model.evaluate(test_dataset)

vgg_loss, vgg_accuracy = vgg_model.evaluate(test_dataset)
print("MobileNetV2 Float32 Accuracy:", mobilenet_accuracy)
print("VGG16 Float32 Accuracy:", vgg_accuracy)

6/6 [=====] - 0s 34ms/step - loss: 0.2244 -
accuracy: 0.9624
```


```
6/6 [=====] - 1s 80ms/step - loss: 0.0947 -  
accuracy: 0.9839  
MobileNetV2 Float32 Accuracy: 0.9623655676841736  
VGG16 Float32 Accuracy: 0.9838709831237793
```

```
import pandas as pd  
  
data = {  
    "Model": ["MobileNetV2", "MobileNetV2", "VGG16", "VGG16"],  
    "Type": ["Float32", "INT8", "Float32", "INT8"],  
    "Accuracy": [  
mobilenet_accuracy,  
mobilenet_int8_acc,  
vgg_accuracy,  
vgg_int8_acc  
],  
    "Model Size (MB)": [  
mobilenet_float_size,  
mobilenet_int8_size,  
vgg_float_size,  
vgg_int8_size  
]  
}  
  
df = pd.DataFrame(data)  
  
df  
  
{  
  "summary": {  
    "name": "df",  
    "rows": 4,  
    "fields": [  
      {  
        "column": "Model",  
        "dtype": "string",  
        "num_unique_values": 2,  
        "samples": ["VGG16", "MobileNetV2"],  
        "semantic_type": "",  
        "description": ""  
      },  
      {  
        "column": "Type",  
        "dtype": "string",  
        "num_unique_values": 2,  
        "samples": ["INT8", "Float32"]  
      }  
    ]  
  }  
}
```



```
plt.ylabel("Model Size (MB)")
plt.title("Model Size Comparison Before and After Quantization")

plt.show()
```




```
models = ["MobileNet Float", "MobileNet INT8", "VGG16 Float", "VGG16 INT8"]

accuracy = [
mobilenet_accuracy,
mobilenet_int8_acc,
vgg_accuracy,
vgg_int8_acc
]

plt.figure(figsize=(8,5))
plt.bar(models, accuracy)

plt.ylabel("Accuracy")
plt.title("Accuracy Comparison Before and After Quantization")

plt.show()
```



Observations

Observations — Post-Training Integer Quantization (INT8)

1. Model Size Reduction

Post-training integer quantization was applied to both MobileNetV2 and VGG16 using TensorFlow Lite, converting 32-bit float weights and activations into 8-bit integers.

- MobileNetV2: 8.97 MB → 2.58 MB — compression ratio of **3.48×**
- VGG16: 57.72 MB → 14.29 MB — compression ratio of **4.04×**

VGG16 achieves a slightly higher compression ratio because its architecture consists mostly of large regular convolutional layers which quantize more uniformly than MobileNetV2's depthwise separable convolutions.

2. Accuracy After Quantization

- MobileNetV2 Float32: **96.24%** → MobileNetV2 INT8: **80.11%** — drop of **~16.13%**
- VGG16 Float32: **98.39%** → VGG16 INT8: **98.39%** — drop of **0%**

VGG16 shows no accuracy loss after INT8 quantization, indicating that its weight distributions are well-suited for linear quantization. MobileNetV2 experiences a significant accuracy drop,

which is common for lightweight models with depthwise separable convolutions — these layers are more sensitive to quantization noise due to their smaller weight ranges per channel.

3. Inference Time

- MobileNetV2 INT8: **0.0043 seconds**
- VGG16 INT8: **0.1309 seconds**

MobileNetV2 is approximately **30× faster** than VGG16 during TFLite inference. This is the key advantage of MobileNetV2 for edge deployment — even with the accuracy drop, it runs extremely fast on CPU-based inference.

4. Summary Table

Model	Type	Accuracy	Size (MB)	Compression Ratio
MobileNetV2	Float32	96.24%	8.97 MB	—
MobileNetV2	INT8	80.11%	2.58 MB	3.48×
VGG16	Float32	98.39%	57.72 MB	—
VGG16	INT8	98.39%	14.29 MB	4.04×

5. Overall Conclusion

INT8 quantization successfully reduces model size by 3.5–4× for both models. VGG16 retains 100% of its accuracy while achieving a 4× compression, making it an excellent candidate for storage-constrained edge deployment. MobileNetV2, despite its accuracy drop, remains far superior in inference speed at just 4ms per sample, making it the better choice for real-time or latency-sensitive edge applications.

Task 4

```
import tensorflow_model_optimization as tfmot
import numpy as np
import pandas as pd
import os

dense_layer = mobilenet_model.layers[-1]

original_weights = dense_layer.get_weights()[0]
bias_weights = dense_layer.get_weights()[1]

print("Dense weight shape:", original_weights.shape)

Dense weight shape: (1280, 3)

from sklearn.cluster import KMeans

num_clusters = 8
```

```

weights_flat = original_weights.reshape(-1,1)
kmeans = KMeans(n_clusters=num_clusters, random_state=0)
kmeans.fit(weights_flat)
cluster_centers = kmeans.cluster_centers_
labels = kmeans.labels_

clustered_weights =
cluster_centers[labels].reshape(original_weights.shape)

quantized_dense = tf.keras.layers.Dense(3, activation="softmax")
quantized_dense.build((None,1280))
quantized_dense.set_weights([clustered_weights,bias_weights])
# Apply clustered weights back to the actual model
dense_layer.set_weights([clustered_weights, bias_weights])
print("Clustered weights applied to MobileNetV2 Dense layer")
Clustered weights applied to MobileNetV2 Dense layer

float_params = original_weights.size * 32
cluster_params = num_clusters * 32 + labels.size *
np.log2(num_clusters)

compression_ratio = float_params / cluster_params
print("Compression Ratio:",compression_ratio)
Compression Ratio: 10.434782608695652

quant_error = np.mean((original_weights - clustered_weights)**2)
print("Quantization Error (MSE):",quant_error)
Quantization Error (MSE): 4.7618585e-05

layer_error_data = {
    "Layer":["Dense"],
    "Quantization Error":[quant_error]
}

df_layer_error = pd.DataFrame(layer_error_data)
df_layer_error

{"summary":{"\n  \"name\": \"df_layer_error\", \n  \"rows\": 1, \n  \"fields\": [\n    {\n      \"column\": \"Layer\", \n

```

```

\ "properties\ ": { \n          \ "dtype\ ": \ "string\ ", \n
\ "num_unique_values\ ": 1, \n          \ "samples\ ": [ \n
\ "Dense\ " \n          ], \n          \ "semantic_type\ ": \ "\", \n
\ "description\ ": \ "\", \n          } \n          }, \n          { \n          \ "column\ ":
\ "Quantization Error\ ", \n          \ "properties\ ": { \n          \ "dtype\ ":
\ "float32\ ", \n          \ "num_unique_values\ ": 1, \n          \ "samples\ ":
[ \n          4.761858508572914e-05 \n          ], \n
\ "semantic_type\ ": \ "\", \n          \ "description\ ": \ "\", \n          } \n
n          } \n          ] \n }", "type": "dataframe", "variable_name": "df_layer_error"

```

```

mobilenet_loss_clustered, mobilenet_acc_clustered =
mobilenet_model.evaluate(test_dataset)
print("Float32 MobileNetV2 Accuracy:", mobilenet_accuracy)
print("Clustered MobileNetV2 Accuracy:", mobilenet_acc_clustered)

```

```

mobilenet_comparison = pd.DataFrame({
    "Model": ["Float32 MobileNetV2", "Clustered MobileNetV2"],
    "Accuracy": [mobilenet_accuracy, mobilenet_acc_clustered]
})
mobilenet_comparison

```

```

6/6 [=====] - 0s 34ms/step - loss: 0.2490 -
accuracy: 0.9677
Float32 MobileNetV2 Accuracy: 0.9623655676841736
Clustered MobileNetV2 Accuracy: 0.9677419066429138

```

```

{ "summary": { \n          \ "name\ ": \ "mobilenet_comparison\ ", \n          \ "rows\ ": 2, \n
\ "fields\ ": [ \n          { \n          \ "column\ ": \ "Model\ ", \n
\ "properties\ ": { \n          \ "dtype\ ": \ "string\ ", \n
\ "num_unique_values\ ": 2, \n          \ "samples\ ": [ \n
\ "Clustered MobileNetV2\ ", \n          \ "Float32 MobileNetV2\ " \n
], \n          \ "semantic_type\ ": \ "\", \n          \ "description\ ": \ "\", \n
} \n          }, \n          { \n          \ "column\ ": \ "Accuracy\ ", \n
\ "properties\ ": { \n          \ "dtype\ ": \ "number\ ", \n          \ "std\ ":
0.0038016457356826417, \n          \ "min\ ": 0.9623655676841736, \n
\ "max\ ": 0.9677419066429138, \n          \ "num_unique_values\ ": 2, \n
\ "samples\ ": [ \n          0.9677419066429138, \n
0.9623655676841736 \n          ], \n          \ "semantic_type\ ": \ "\", \n
\ "description\ ": \ "\", \n          } \n          } \n          ] \n
n }", "type": "dataframe", "variable_name": "mobilenet_comparison"

```

```

import matplotlib.pyplot as plt

```

```

plt.hist(original_weights.flatten(), bins=50,
alpha=0.5, label="Original")
plt.hist(clustered_weights.flatten(), bins=50,
alpha=0.5, label="Clustered")


```

```

plt.title("Weight Distribution Before and After Clustering")
plt.legend()

```

```
plt.show()
```



```
print("Mean Quantization Error:", quant_error)
print("Maximum Error:",
      np.max(np.abs(original_weights - clustered_weights)))
print("Minimum Error:",
      np.min(np.abs(original_weights - clustered_weights)))
```

Mean Quantization Error: 4.7618585e-05

Maximum Error: 0.034852363

Minimum Error: 2.3003668e-06

```
import matplotlib.pyplot as plt
```


```
weight_diff = original_weights - clustered_weights
```

```
plt.hist(weight_diff.flatten(), bins=50)
```

```
plt.title("Weight Difference Distribution")
```

```
plt.xlabel("Weight Error")
```

```
plt.ylabel("Frequency")
plt.show()
```



```
print("Cluster Centers:")
print(cluster_centers)

Cluster Centers:
[[ 0.00103973]
 [-0.03842615]
 [ 0.04334426]
 [-0.0185078 ]
 [ 0.0704391 ]
 [-0.0840314 ]
 [ 0.02155362]
 [-0.05890602]]

unique_original = len(np.unique(original_weights))
unique_clustered = len(np.unique(clustered_weights))

print("Unique weights (Original):", unique_original)
```

```
print("Unique weights (Clustered):", unique_clustered)

Unique weights (Original): 3840
Unique weights (Clustered): 8


plt.hist(labels, bins=num_clusters)

plt.title("Cluster Assignment Distribution")

plt.xlabel("Cluster Index")

plt.ylabel("Number of Weights")

plt.show()
```




```
plt.scatter(range(len(cluster_centers)), cluster_centers)

plt.title("Cluster Centroids")

plt.xlabel("Cluster ID")

plt.ylabel("Weight Value")
```

```
plt.show()
```



For vgg16

```
import numpy as np
import pandas as pd
from sklearn.cluster import KMeans
import tensorflow as tf
import matplotlib.pyplot as plt

n_bits = 3
num_clusters = 2 ** n_bits

print("Number of Clusters:", num_clusters)

Number of Clusters: 8

dense_layer = vgg_model.layers[-1]

weights = dense_layer.get_weights()

kernel = weights[0]
bias = weights[1]
```

```

flat_weights = kernel.reshape(-1,1)

kmeans = KMeans(n_clusters=num_clusters, random_state=0)

labels = kmeans.fit_predict(flat_weights)

clustered_weights =
kmeans.cluster_centers_[labels].reshape(kernel.shape)

quant_error = np.mean((kernel - clustered_weights) ** 2)

print("Quantization Error:", quant_error)
Quantization Error: 0.00039645875

dense_layer.set_weights([clustered_weights, bias])

original_unique = len(np.unique(kernel))
cluster_unique = len(np.unique(clustered_weights))

compression_ratio = original_unique / cluster_unique

print("Original unique weights:", original_unique)
print("Clustered unique weights:", cluster_unique)
print("Compression Ratio:", compression_ratio)

Original unique weights: 768
Clustered unique weights: 8
Compression Ratio: 96.0

loss, acc = vgg_model.evaluate(test_dataset_vgg)

print("Clustered VGG16 Accuracy:", acc)


12/12 [=====] - 1s 46ms/step - loss: 0.0943 -
accuracy: 0.9839
Clustered VGG16 Accuracy: 0.9838709831237793

plt.hist(kernel.flatten(), bins=50)
plt.title("Original Weight Distribution")
plt.show()


plt.hist(clustered_weights.flatten(), bins=50)
plt.title("Clustered Weight Distribution")
plt.show()

```


Original Weight Distribution



Clustered Weight Distribution



```

df_layer_errors = pd.DataFrame(
    [{"Dense Layer", quant_error}],
    columns=["Layer", "Quantization Error"]
)

df_layer_errors

{"summary": "{\n  \"name\": \"df_layer_errors\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Layer\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 1,\n        \"samples\": [\n          \"Dense Layer\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Quantization Error\",\n      \"properties\": {\n        \"dtype\": \"float32\",\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.00039645875222049654\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"df_layer_errors\"}"}

loss, acc = vgg_model.evaluate(test_dataset)

print("Clustered VGG16 Accuracy:", acc)

6/6 [=====] - 1s 83ms/step - loss: 0.0943 - accuracy: 0.9839
Clustered VGG16 Accuracy: 0.9838709831237793

comparison = pd.DataFrame({
    "Model": ["Float32 VGG16", "Clustered VGG16"],
    "Accuracy": [vgg_accuracy, acc]
})


comparison

{"summary": "{\n  \"name\": \"comparison\",\n  \"rows\": 2,\n  \"fields\": [\n    {\n      \"column\": \"Model\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"Clustered VGG16\", \"Float32 VGG16\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.0,\n        \"min\": 0.9838709831237793,\n        \"max\": 0.9838709831237793,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.9838709831237793\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"comparison\"}"}

labels = ["Original", "Clustered"]
values = [original_unique, cluster_unique]

```


```
plt.bar(labels, values)
plt.title("Unique Weights Reduction After Clustering")
plt.ylabel("Number of Unique Weights")
plt.show()
```



```
plt.bar(["Dense Layer"], [quant_error])
plt.title("Quantization Error After Clustering")
plt.ylabel("MSE Error")
plt.show()
```



```
plt.figure(figsize=(10,4))  
  
plt.subplot(1,2,1)  
plt.hist(kernel.flatten(), bins=50)  
plt.title("Original Weights")  
  
plt.subplot(1,2,2)  
plt.hist(clustered_weights.flatten(), bins=50)  
plt.title("Clustered Weights")  
  
plt.show()
```



```
print("Clusters Used:", num_clusters)
print("Compression Ratio:", compression_ratio)
print("Quantization Error:", quant_error)
print("Clustered Accuracy:", acc)
```

```
Clusters Used: 8
Compression Ratio: 96.0
Quantization Error: 0.00039645875
Clustered Accuracy: 0.9838709831237793
```

Observations — n-bit Clustering Quantization

1. Clustering Configuration

3-bit clustering quantization was applied to both models, resulting in **8 cluster centroids ($2^3 = 8$)** used to represent all weight values in the final Dense layer. The K-Means algorithm grouped similar weights and replaced each original weight with its nearest cluster centroid.

2. Compression Ratio

- **MobileNetV2:** Original unique weights: 3840 → Clustered: 8 — Compression Ratio: **10.43×**
- **VGG16:** Original unique weights: 768 → Clustered: 8 — Compression Ratio: **96.0×**

VGG16 achieves a dramatically higher compression ratio because its Dense layer has fewer unique weight values relative to the number of clusters, resulting in near-complete consolidation. This demonstrates that clustering-based quantization is especially effective for layers with redundant weight distributions.

3. Quantization Error (MSE)

- **MobileNetV2 Dense Layer MSE:** 4.76×10^{-5}
- **VGG16 Dense Layer MSE:** 3.96×10^{-4}

MobileNetV2 has a lower quantization error, indicating its Dense layer weights are more tightly clustered around centroid values. VGG16 has a slightly higher error due to its wider weight distribution (ranging from -0.3 to +0.3 as seen in the histogram), which makes it harder to represent all values with just 8 centroids.

4. Weight Distribution Analysis

From the weight distribution plots:

MobileNetV2: The original weights follow a roughly bell-shaped distribution centered near zero, spanning approximately -0.10 to +0.10. After clustering, the weights collapse into 8 discrete spikes at the centroid positions. The weight difference distribution is tightly concentrated around zero (within ± 0.015), confirming low reconstruction error.

VGG16: The original weights span a wider range (-0.3 to +0.3) with a relatively uniform distribution. After clustering, 8 distinct centroid bars replace the continuous distribution. The unique weights reduction bar chart clearly shows the dramatic drop from 768 $\rightarrow$ 8 unique values.

5. Cluster Centroid Analysis

The 8 cluster centroids for MobileNetV2 are: [-0.084, -0.059, -0.038, -0.019, +0.001, +0.022, +0.043, +0.070]

The centroids are evenly spread across the weight range, and the cluster assignment distribution shows that clusters 0 and 6 contain the most weights (620 each), while cluster 5 contains the fewest (160), indicating that weights are not uniformly distributed across the value range.

6. Model Accuracy After Clustering

Model	Float32 Accuracy	Clustered Accuracy	Change
MobileNetV2	96.24%	96.77%	+0.54%
VGG16	98.39%	98.39%	0.00%

Remarkably, both models maintain or slightly improve their accuracy after clustering. MobileNetV2 even shows a minor accuracy gain of +0.54%, which can be attributed to a slight regularization effect from weight quantization. VGG16 shows zero accuracy degradation despite the extreme 96 $\times$ compression ratio.

7. Overall Conclusion

Clustering-based quantization applied to the Dense layer achieves significant compression (10 $\times$ for MobileNetV2, 96 $\times$ for VGG16) with negligible quantization error and no accuracy loss. This demonstrates that the final classification layer of both models contains highly redundant weight values that can be effectively represented by just 8 discrete centroids. The technique is highly suitable for edge deployment where storage of model weights is a constraint.

Task 5

```
import numpy as np
import pandas as pd

def linear_quantize(weights, bits):

    qmin = -(2**(bits-1))
    qmax = (2**(bits-1)) - 1

    w_min = weights.min()
    w_max = weights.max()

    scale = (w_max - w_min) / (qmax - qmin)

    zero_point = qmin - w_min / scale

    q_weights = np.round(weights / scale + zero_point)

    q_weights = np.clip(q_weights, qmin, qmax)

    reconstructed = (q_weights - zero_point) * scale

    return reconstructed

mobilenet_int8_model = tf.keras.models.clone_model(mobilenet_model)
mobilenet_int8_model.set_weights(mobilenet_model.get_weights())

layer_errors = []

for layer in mobilenet_int8_model.layers:

    weights = layer.get_weights()

    if len(weights) == 0:
        continue

    new_weights = []

    for w in weights:

        quantized = linear_quantize(w, 8)

        error = np.mean((w - quantized)**2)

        layer_errors.append((layer.name, error))

        new_weights.append(quantized)

    layer.set_weights(new_weights)
```

```

df_int8_mobilenet =
pd.DataFrame(layer_errors,columns=["Layer","Quantization Error"])
df_int8_mobilenet

{"summary":{"\n  \"name\": \"df_int8_mobilenet\", \n  \"rows\": 262, \n
\"fields\": [\n    {\n      \"column\": \"Layer\", \n
\"properties\": {\n        \"dtype\": \"category\", \n
\"num_unique_values\": 2, \n        \"samples\": [\n
\"dense\", \n        \"mobilenetv2_1.00_128\", \n        ], \n
\"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n
    }, \n    {\n      \"column\": \"Quantization Error\", \n
\"properties\": {\n        \"dtype\": \"float32\", \n
\"num_unique_values\": 262, \n        \"samples\": [\n
0.00031687895534560084, \n        4.466715836315416e-05 \n        ], \n
\"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n
    ] \n  } \n}, \"type\": \"dataframe\", \"variable_name\": \"df_int8_mobilenet\"}

mobilenet_int16_model = tf.keras.models.clone_model(mobilenet_model)
mobilenet_int16_model.set_weights(mobilenet_model.get_weights())

layer_errors = []

for layer in mobilenet_int16_model.layers:

    weights = layer.get_weights()

    if len(weights) == 0:
        continue

    new_weights = []

    for w in weights:

        quantized = linear_quantize(w,16)

        error = np.mean((w - quantized)**2)

        layer_errors.append((layer.name,error))

        new_weights.append(quantized)

    layer.set_weights(new_weights)

df_int16_mobilenet =
pd.DataFrame(layer_errors,columns=["Layer","Quantization Error"])
df_int16_mobilenet

{"summary":{"\n  \"name\": \"df_int16_mobilenet\", \n  \"rows\": 262, \n
\"fields\": [\n    {\n      \"column\": \"Layer\", \n
\"properties\": {\n        \"dtype\": \"category\", \n

```



```

\"num_unique_values\": 2,\n          \"samples\": [\n
\"dense\",,\n          \"mobilenetv2_1.00_128\",,\n
\"semantic_type\": \"\",,\n          \"description\": \"\",,\n
n    },\n    {\n          \"column\": \"Quantization Error\",,\n
\"properties\": {\n          \"dtype\": \"float32\",,\n
\"num_unique_values\": 262,\n          \"samples\": [\n
4.79985562407137e-09,\n          6.755133674651859e-10\n          ],,\n
\"semantic_type\": \"\",,\n          \"description\": \"\",,\n
n    },\n  ]\n
n}","type":"dataframe","variable_name":"df_int16_mobilenet"}

```

for vgg16

```

vgg_int8_model = tf.keras.models.clone_model(vgg_model)
vgg_int8_model.set_weights(vgg_model.get_weights())

layer_errors = []

for layer in vgg_int8_model.layers:
    weights = layer.get_weights()

    if len(weights) == 0:
        continue

    new_weights = []

    for w in weights:
        quantized = linear_quantize(w,8)

        error = np.mean((w - quantized)**2)

        layer_errors.append((layer.name,error))

        new_weights.append(quantized)

    layer.set_weights(new_weights)

df_int8_vgg = pd.DataFrame(layer_errors,columns=["Layer","Quantization
Error"])
df_int8_vgg

{"summary":{"\n  \"name\": \"df_int8_vgg\",,\n  \"rows\": 30,\n
\"fields\": [\n    {\n          \"column\": \"Layer\",,\n
\"properties\": {\n          \"dtype\": \"category\",,\n
\"num_unique_values\": 3,\n          \"samples\": [\n
\"vgg16\",,\n          \"dense\",,\n          \"dense_1\",,\n
\"semantic_type\": \"\",,\n          \"description\": \"\",,\n
n    },\n    {\n          \"column\": \"Quantization Error\",,\n
\"properties\": {\n          \"dtype\": \"float32\",,\n

```

```

{"num_unique_values": 30,\n
  "samples": [\n
    1.792906267894523e-09,\n
    2.751428098690667e-07,\n
    3.7347608667914756e-06\n
  ],\n
  "semantic_type": "\"",\n
  "description": "\"\""}\n
n}","type":"dataframe","variable_name":"df_int8_vgg"}

```

```
vgg_int16_model = tf.keras.models.clone_model(vgg_model)
vgg_int16_model.set_weights(vgg_model.get_weights())
```

```
layer_errors = []
```

```
for layer in vgg_int16_model.layers:
```

```
weights = layer.get_weights()
```

```
if len(weights) == 0:
    continue
```

```
new_weights = []
```

```
for w in weights:
```

```
quantized = linear_quantize(w, 16)
```

```
error = np.mean((w - quantized)**2)
```

```
layer_errors.append((layer.name,error))
```

```
new_weights.append(quantized)
```

```
layer.set_weights(new_weights)
```

```
df_int16_vgg =
pd.DataFrame(layer_errors, columns=["Layer", "Quantization Error"])
df_int16_vgg
```

```
{
  "summary": {
    "name": "df_int16_vgg",
    "rows": 30,
    "fields": [
      {
        "column": "Layer",
        "dtype": "category",
        "num_unique_values": 3,
        "samples": [
          "vgg16",
          "dense",
          "dense_1"
        ],
        "semantic_type": "",
        "description": ""
      },
      {
        "column": "Quantization Error",
        "dtype": "float32",
        "num_unique_values": 30,
        "samples": [
          2.978816330970538e-14,
          4.1443059921320025e-12,
          5.236102523786812e-11
        ],
        "semantic_type": "",
        "description": ""
      }
    ]
  },
  "type": "dataframe",
  "variable name": "df_int16_vgg"
}
```

```

def model_error(original,quantized):

    errors = []

    for w1,w2 in zip(original.get_weights(),quantized.get_weights()):
        errors.append(np.mean((w1-w2)**2))

    return np.mean(errors)

print("MobileNet INT8
Error:",model_error(mobilenet_model,mobilenet_int8_model))
print("MobileNet INT16
Error:",model_error(mobilenet_model,mobilenet_int16_model))

print("VGG16 INT8 Error:",model_error(vgg_model,vgg_int8_model))
print("VGG16 INT16 Error:",model_error(vgg_model,vgg_int16_model))

MobileNet INT8 Error: 0.0009953339
MobileNet INT16 Error: 1.5578964e-08
VGG16 INT8 Error: 4.9325126e-06
VGG16 INT16 Error: 7.996015e-11

mobilenet_int8_model.compile(optimizer="adam",
loss="sparse_categorical_crossentropy", metrics=["accuracy"])
mobilenet_int16_model.compile(optimizer="adam",
loss="sparse_categorical_crossentropy", metrics=["accuracy"])
vgg_int8_model.compile(optimizer="adam",
loss="sparse_categorical_crossentropy", metrics=["accuracy"])
vgg_int16_model.compile(optimizer="adam",
loss="sparse_categorical_crossentropy", metrics=["accuracy"])

print("MobileNet INT8 Accuracy")
mobilenet_int8_model.evaluate(test_dataset)

print("MobileNet INT16 Accuracy")
mobilenet_int16_model.evaluate(test_dataset)

print("VGG16 INT8 Accuracy")
vgg_int8_model.evaluate(test_dataset)

print("VGG16 INT16 Accuracy")
vgg_int16_model.evaluate(test_dataset_vgg)

MobileNet INT8 Accuracy
6/6 [=====] - 2s 38ms/step - loss: 0.2373 -
accuracy: 0.9409
MobileNet INT16 Accuracy
6/6 [=====] - 2s 43ms/step - loss: 0.2496 -
accuracy: 0.9624
VGG16 INT8 Accuracy

```

```
6/6 [=====] - 1s 82ms/step - loss: 0.0963 -  
accuracy: 0.9839  
VGG16 INT16 Accuracy  
12/12 [=====] - 1s 41ms/step - loss: 0.0943 -  
accuracy: 0.9839
```

```
[0.09427588433027267, 0.9838709831237793]
```

```
def compression_ratio(original_bits, quant_bits):  
    return original_bits / quant_bits
```

```
print("INT8 Compression Ratio:", compression_ratio(32,8))  
print("INT16 Compression Ratio:", compression_ratio(32,16))
```

```
INT8 Compression Ratio: 4.0  
INT16 Compression Ratio: 2.0
```

```
results = pd.DataFrame({  
    "Model": [  
        "MobileNet Float32",  
        "MobileNet INT8",  
        "MobileNet INT16",  
        "VGG16 Float32",  
        "VGG16 INT8",  
        "VGG16 INT16"  
    ],  
    "Accuracy": [  
        mobilenet_accuracy,  
        mobilenet_int8_model.evaluate(test_dataset, verbose=0)[1],  
        mobilenet_int16_model.evaluate(test_dataset, verbose=0)[1],  
        vgg_accuracy,  
        vgg_int8_model.evaluate(test_dataset_vgg, verbose=0)[1],  
        vgg_int16_model.evaluate(test_dataset_vgg, verbose=0)[1]  
    ]  
})  
results
```


```
{"summary": "{\n  \"name\": \"results\",\n  \"rows\": 6,\n  \"fields\": [\n    {\n      \"column\": \"Model\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 6,\n        \"samples\": [\n          \"MobileNet Float32\",\n          \"MobileNet INT8\",\n          \"VGG16 INT16\"  
        ],\n        \"semantic_type\": \"\",  
        ]\n      }\n    }\n  ]\n}
```

```

\"semantic_type\": \"\", \n      \"description\": \"\" \n    } \n  ] \n} \", \"type\": \"dataframe\", \"variable_name\": \"results\"}

plt.figure(figsize=(10,5))
plt.bar(results[\"Model\"],results[\"Accuracy\"] )
plt.xticks(rotation=45)
plt.title(\"Accuracy Comparison (Float vs INT8 vs INT16)\")
plt.ylabel(\"Accuracy\")
plt.show()


```



```

plt.figure(figsize=(10,5))
plt.bar(df_int8_vgg[\"Layer\"],df_int8_vgg[\"Quantization Error\"] )
plt.xticks(rotation=90)
plt.title(\"Layer-wise Quantization Error (INT8 VGG16)\")
plt.show()

```




```
weights = vgg_model.layers[-1].get_weights()[0]
quantized_weights = vgg_int8_model.layers[-1].get_weights()[0]
plt.figure(figsize=(10,4))

plt.subplot(1,2,1)
plt.hist(weights.flatten(),bins=50)
plt.title("Original Weights")



plt.subplot(1,2,2)
plt.hist(quantized_weights.flatten(),bins=50)
plt.title("INT8 Quantized Weights")

plt.show()
```




```
# Layer-wise Quantization Error - MobileNetV2 INT8
plt.figure(figsize=(10,4))
plt.bar(df_int8_mobilenet["Layer"], df_int8_mobilenet["Quantization
Error"])
plt.xticks(rotation=90)
plt.title("Layer-wise Quantization Error (INT8 MobileNetV2)")
plt.ylabel("MSE Error")
plt.tight_layout()
plt.show()

# Layer-wise Quantization Error - MobileNetV2 INT16
plt.figure(figsize=(10,4))
plt.bar(df_int16_mobilenet["Layer"], df_int16_mobilenet["Quantization
Error"])
plt.xticks(rotation=90)
plt.title("Layer-wise Quantization Error (INT16 MobileNetV2)")
plt.ylabel("MSE Error")
plt.tight_layout()
plt.show()
```



```
plt.figure(figsize=(10,4))
plt.bar(df_int16_vgg["Layer"], df_int16_vgg["Quantization Error"])
plt.xticks(rotation=90)
plt.title("Layer-wise Quantization Error (INT16 VGG16)")
plt.ylabel("MSE Error")
plt.tight_layout()
plt.show()
```



```
# Weight distribution before vs after INT8 for MobileNetV2
weights_mob = mobilenet_model.layers[-1].get_weights()[0]
quantized_weights_mob = mobilenet_int8_model.layers[-1].get_weights()
[0]

plt.figure(figsize=(10,4))

plt.subplot(1,2,1)
plt.hist(weights_mob.flatten(), bins=50)
plt.title("MobileNetV2 Original Weights")

plt.subplot(1,2,2)
plt.hist(quantized_weights_mob.flatten(), bins=50)
plt.title("MobileNetV2 INT8 Quantized Weights")

plt.tight_layout()
plt.show()
```



```

summary = pd.DataFrame({
    "Model": [
        "MobileNet Float32",
        "MobileNet INT8",
        "MobileNet INT16",
        "VGG16 Float32",
        "VGG16 INT8",
        "VGG16 INT16"
    ],
    "Bits": [32, 8, 16, 32, 8, 16],
    "Compression Ratio": ["1x", "4x", "2x", "1x", "4x", "2x"],
    "Whole Model MSE Error": [
        0,
        model_error(mobilenet_model, mobilenet_int8_model),
        model_error(mobilenet_model, mobilenet_int16_model),
        0,
        model_error(vgg_model, vgg_int8_model),
        model_error(vgg_model, vgg_int16_model)
    ]
})

```

summary

```

{"summary":{"\n  \"name\": \"summary\",\n  \"rows\": 6,\n  \"fields\": [\n    {\n      \"column\": \"Model\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 6,\n        \"samples\": [\n          \"MobileNet Float32\",\n          \"MobileNet INT8\",\n          \"VGG16 INT16\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Bits\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 10,\n        \"min\": 8,\n        \"max\": 32,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          32,\n          8,\n          16\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Compression Ratio\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 3,\n        \"samples\": [\n          \"1\\u00d7\",\n          \"4\\u00d7\",\n          \"2\\u00d7\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Whole Model MSE Error\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.00040594412908323427,\n        \"min\": 0.0,\n        \"max\": 0.000995333888567984,\n        \"num_unique_values\": 5,\n        \"samples\": [\n          0.000995333888567984,\n          7.996014961264564e-11,\n          1.557896389670077e-08\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n},\"type\":\"dataframe\",\"variable_name\":\"summary"}

```


```

labels = ["MobileNet INT8", "MobileNet INT16", "VGG16 INT8", "VGG16 INT16"]

errors = [
    model_error(mobilenet_model, mobilenet_int8_model),
    model_error(mobilenet_model, mobilenet_int16_model),
    model_error(vgg_model, vgg_int8_model),
    model_error(vgg_model, vgg_int16_model)
]

plt.figure(figsize=(8,5))
plt.bar(labels, errors)
plt.title("Whole Model Quantization Error Comparison")
plt.ylabel("Mean Squared Error")
plt.xticks(rotation=15)
plt.tight_layout()
plt.show()

```



```

print("Linear Quantization Summary")
print("-----")
print("INT8 Compression Ratio: 4x")
print("INT16 Compression Ratio: 2x")
print("INT8 Quantization introduces higher error but better

```

```
compression")
print("INT16 Quantization preserves higher precision")

Linear Quantization Summary
-----
INT8 Compression Ratio: 4x
INT16 Compression Ratio: 2x
INT8 Quantization introduces higher error but better compression
INT16 Quantization preserves higher precision
```

Observations — Linear Quantization (INT8 and INT16)

1. Quantization Method

Linear quantization was applied to convert float32 weights into lower-precision integer representations (INT8 and INT16) using a min-max affine quantization scheme with a computed scale factor and zero-point for each weight tensor.

2. Compression Ratios

- **INT8 Compression Ratio: 4×** (32-bit → 8-bit)
- **INT16 Compression Ratio: 2×** (32-bit → 16-bit)

These ratios are fixed and apply equally to both models since compression depends only on the bit-width reduction, not the model architecture.

3. Whole Model Quantization Error (MSE)

Model	INT8 Error	INT16 Error
MobileNetV2	9.95×10^{-4}	1.56×10^{-8}
VGG16	4.93×10^{-6}	7.99×10^{-11}

INT8 quantization introduces significantly higher error than INT16 for both models. MobileNetV2 has a notably higher INT8 error (~ 0.001) compared to VGG16 ($\sim 4.93 \times 10^{-6}$), which is because MobileNetV2's depthwise separable convolution weights have a wider dynamic range that is harder to represent with only 256 discrete levels. INT16 error is extremely small for both models, confirming near-lossless weight reconstruction at 16-bit precision.

4. Layer-wise Quantization Error Analysis

MobileNetV2: The layer-wise error plots show that the base MobileNetV2 convolutional layers contribute the overwhelming majority of the INT8 quantization error (0.09 MSE), while the Dense classifier layer has negligible error ($\sim 3.5 \times 10^{-8}$). This confirms that the depthwise separable convolutions in the base model are the most sensitive to INT8 quantization. For INT16, all layer errors drop to the order of 10^{-10} to 10^{-6} , confirming excellent weight preservation.

VGG16: The INT8 layer-wise plot shows that the VGG16 base convolutional layers dominate the error (max 1.2×10^{-4}), while the Dense layers have very small errors (5×10^{-7}). Notably, layer index 25 (the last convolutional layer) has the highest error of all VGG16 layers. For INT16, all errors drop to the range of 10^{-12} to 10^{-9} , essentially lossless.

5. Weight Distribution Analysis

The weight histogram comparisons confirm that:

MobileNetV2: The original Dense layer weights (clustered from Task 4) show 8 discrete spikes. After INT8 quantization, the histogram is nearly identical — the discrete nature of already-clustered weights means INT8 introduces minimal additional distortion.

VGG16: The original Dense layer weights also show 8 discrete spikes from Task 4 clustering. After INT8 quantization, the histogram remains visually identical, confirming very low quantization error for the Dense layer specifically.

6. Model Accuracy After Linear Quantization

Model	Float32	INT8	INT16
MobileNetV2	96.24%	94.09%	96.24%
VGG16	98.39%	98.39%	98.39%

MobileNetV2 INT8 experiences a moderate accuracy drop of ~2.15%, consistent with its higher quantization error. **MobileNetV2 INT16** perfectly preserves the original accuracy at 96.24%, demonstrating that 16-bit precision is sufficient for lossless quantization of this model.

VGG16 shows zero accuracy degradation for both INT8 and INT16, maintaining 98.39% across all precision levels. This remarkable result confirms that VGG16's regular convolutional weight distributions are highly amenable to linear quantization.

7. Overall Conclusion

Linear quantization is highly effective for model compression with minimal accuracy trade-off. INT16 achieves 2× compression with essentially zero accuracy loss for both models, making it the recommended choice when accuracy preservation is critical. INT8 achieves 4× compression — VGG16 handles this perfectly with no accuracy loss, while MobileNetV2 shows a small ~2% drop which may be acceptable for edge deployment where memory constraints are more important than marginal accuracy differences. In all cases, the quantization error is extremely small relative to the weight magnitudes, confirming that linear quantization is a practical and efficient technique for deploying deep learning models on resource-constrained edge devices.

Overall Assignment Conclusion

Summary of All Tasks

This assignment explored model compression techniques for edge AI deployment using MobileNetV2 and VGG16 trained on the Rock-Paper-Scissors dataset.

Task 1 — Model Training

Both models were successfully fine-tuned with the top 10 layers unfrozen:

- **MobileNetV2** achieved **96.24%** test accuracy (lightweight, 2.26M parameters)

- **VGG16** achieved **98.39%** test accuracy (heavyweight, 14.85M parameters)

Task 2 — Efficiency Metrics

Metric	MobileNetV2	VGG16
Parameters	2.26M	14.85M
Model Size	8.63 MB	56.64 MB
FLOPs	195.69 M	10,028.98 M
Inference Time	0.87s	0.28s

VGG16 is 6.6× larger and 51× more compute-intensive than MobileNetV2, confirming that MobileNetV2 is far more suitable for edge deployment.

Task 3 — TFLite INT8 Post-Training Quantization

Model	Size Reduction	Accuracy Drop
MobileNetV2	8.97 MB → 2.58 MB (3.48×)	96.24% → 80.11% (−16.13%)
VGG16	57.72 MB → 14.29 MB (4.04×)	98.39% → 98.39% (0%)

TFLite quantization is highly effective for VGG16 but causes significant accuracy loss in MobileNetV2, suggesting it requires quantization-aware training for better results.

Task 4 — Clustering Quantization (3-bit, 8 clusters)

Model	Compression Ratio	MSE Error	Accuracy Change
MobileNetV2	10.43×	4.76×10^{-5}	96.24% → 96.77% (+0.54%)
VGG16	96.0×	3.96×10^{-4}	98.39% → 98.39% (0%)

Clustering quantization achieves extreme compression ratios with negligible accuracy loss, demonstrating that both models contain highly redundant weight values in their Dense layers.

Task 5 — Linear Quantization (INT8 and INT16)

Model	INT8 Error	INT16 Error	INT8 Accuracy	INT16 Accuracy
MobileNetV2	9.95×10^{-4}	1.56×10^{-8}	94.09% (−2.15%)	96.24% (0%)
VGG16	4.93×10^{-6}	7.99×10^{-11}	98.39% (0%)	98.39% (0%)

INT16 quantization achieves 2× compression with zero accuracy loss for both models. INT8 achieves 4× compression — perfect for VGG16, with only a minor drop for MobileNetV2.

Key Takeaways

1. **MobileNetV2 is better suited for edge deployment** due to its lightweight architecture (8.63 MB, 195M FLOPs) compared to VGG16 (56.64 MB, 10,028M FLOPs).
2. **VGG16 is more robust to quantization** — it maintains 98.39% accuracy across all compression techniques (TFLite INT8, clustering, linear INT8/INT16), while MobileNetV2 is more sensitive to aggressive quantization.
3. **Best compression technique per use case:**
 - Maximum compression with no accuracy loss → **Clustering (96× for VGG16)**
 - Best accuracy-compression trade-off → **Linear INT16 (2× compression, 0% drop)**
 - Smallest final model size → **TFLite INT8 (2.58 MB for MobileNetV2)**
4. **Convolutional layers are the primary source of quantization error** in both models, while Dense layers are highly resilient to all quantization methods tested.
5. All three quantization techniques confirm that deep learning models can be significantly compressed for edge deployment while retaining competitive accuracy, making them practical for real-world applications on resource-constrained devices.